

# MONITORING & MANAGEMENT WITH DOCKER & CONTAINERS

EDITED & CURATED BY ALEX WILLIAMS

## **The New Stack: The Docker and Container Ecosystem Ebook Series**

Alex Williams, Founder & Editor-in-Chief  
Benjamin Ball, Technical Editor & Producer  
Gabriel Hoang Dinh, Creative Director  
Lawrence Hecht, Data Research Director

### **Contributors:**

Judy Williams, Copy Editor  
Luke Lefler, Audio Engineer  
Norris Deajon, Audio Engineer

# TABLE OF CONTENTS

Sponsors .....	4
Introduction .....	5

## **MONITORING & MANAGEMENT WITH DOCKER & CONTAINERS**


Monitoring Reset for Containers.....	8
IBM: Creating Analytics-Driven Solutions for Operational Visibility.....	23
Classes of Container Monitoring .....	24
Docker: Building On Docker's Native Monitoring Functionality.....	40
Identifying and Collecting Container Data.....	41
Sysdig: The Importance of Having Visibility Into Containers.....	55
The Right Tool for the Job: Picking a Monitoring Solution .....	56
Cisco: Understanding the Application Pattern for Effective Monitoring .....	63
Managing Intentional Chaos: Succeeding with Containerized Apps.....	64
Monitoring Docker and Kubernetes at WayBlazer.....	71

## **CONTAINER MONITORING DIRECTORY**

Container-Related Monitoring.....	80
Components/Classes of Monitoring Systems .....	83
Management/Orchestration .....	87
Miscellaneous .....	89
Disclosures.....	91

# SPONSORS

We are grateful for the support of the following ebook series sponsors:



And the following sponsor for this ebook:



# INTRODUCTION

As we've produced The Docker & Container Ecosystem ebook series, it feels like container usage has become more normalized. We talk less about the reasons why to use and adopt containers, and speak more to the challenges of widespread container deployments. With each book in the series, we've focused more on not just moving container workloads into production environments, but keeping those systems healthy and operational.

With the introduction of containers and microservices, monitoring solutions have to handle more ephemeral services and server instances than ever before. Not only are there more objects to manage, but they're generating more pieces of information. Collecting data from environments composed of so many moving parts has become increasingly complex. Monitoring containerized applications and environments has also moved beyond just the operations team, or at the very least has opened up to more roles, as the DevOps movement encourages cross-team accountability.

This book starts with a look at what monitoring applications and systems has typically meant for teams, and how the environments and goals have changed. Monitoring as a behavior has come a long way from just keeping computers online. While many of the core behaviors remain similar, there's a lot that's different with containers. Many of the defining characteristics of containers are what make them difficult to monitor. We discuss these differences, and how users and vendors are adapting to meet the necessities of containerized environments.

Moving beyond historical context, we learn about the components, behaviors, use cases and event types that compose container monitoring. There is a breakdown of the different approaches for measuring and

handling events, including metrics, logs, profiling and tracing. This breakdown includes a discussion about metrics-based monitoring systems, and how the metrics approach constitutes a series of functions such as collection, ingestion, storage, processing, alerting and visualization.

Another major focus of this book is how users are actually collecting the monitoring data from the containers they're utilizing. For some, the data collection is as simple as the native capabilities built into the Docker engine. Docker's base monitoring capabilities are almost like an agnostic platform for other tools that build off it. We cover some of the different tools for data collection, which range from self-hosted open source solutions to commercial cloud-based services that are built specifically for containers. We also discuss how users and teams can pick a monitoring solution that best fits their requirements.

We cover how containerized applications and services can be designed with consideration for container-native patterns and the operations teams that will manage them. Creating applications that are operations-aware will enable faster deployment and better communication across teams. This kind of focus means creating applications and services that are designed for fast startups, graceful shutdowns and the complexity of container management systems.

Much like the container ecosystem itself, the approaches to monitoring containers are varied and opinionated, but present many valuable opportunities. The monitoring landscape is full of solutions for many different types of environments and teams. These monitoring solutions include many open source and commercial solutions that range from native capabilities built into cloud platforms to cloud-based services purpose-built for containers. We try to cover as many of these monitoring perspectives as possible, including the very real possibilities of using more than one solution.

While this book is the last in its series, this by no means our last publication on the container ecosystem. If anything, our perspective on how to approach this landscape has changed drastically — it seems almost impossible not to try and narrow down the container ecosystem into more approachable communities. If we've learned anything, it's that this ecosystem is incredibly lush and growing at a rate that is difficult to consistently portray.

As we refine our ecosystem perspective, we are looking to break it down into new areas of focus, such as Docker, Kubernetes and serverless technologies. Expect to see more from us on building and managing scale out ecosystems. We've learned so much from this first ebook series on Docker and containers, and we hope you've learned a few things as well. I'd like to thank all our sponsors, writers and readers for how far we've come in the last year, and we'll see you all again soon.


Thanks,

**Benjamin Ball**

Technical Editor and Producer  
The New Stack

# MONITORING RESET FOR CONTAINERS

by **LAWRENCE HECHT**

onitoring is not a new concept, but a lot has changed about the systems that need monitoring and which teams are responsible for it. In the past, monitoring used to be as simple as checking if a computer was still running. Cobe Chief Technical Officer (CTO) Dave Charles [remembers monitoring](#) as simple instrumentation that came alongside a product.

As James Turnbull explains in [The Art of Monitoring](#), most small organizations didn't have automated monitoring — they instead focused on minimizing downtime and managing physical assets. At companies that actually had IT staff, operations teams used simple tools to check on disk, central processing unit (CPU) and memory usage, but focused mostly on dealing with emergencies related to availability. Larger organizations eventually replaced the manual approach with automated monitoring systems that utilized dashboards.

Even without the introduction of containers, recent thought leaders have advocated that monitoring should more proactively look at ways to improve performance. To get a better view of the monitoring environment,

## Analysis of Monitoring Status Quo Circa 2015

Survey Results	Analysis
<p><b>Nagios was the most widely used tool, followed by Amazon CloudWatch and New Relic.</b> A second tier of products included <a href="#">Incinga</a>, <a href="#">Sensu</a>, <a href="#">Zabbix</a>, <a href="#">Datadog</a>, <a href="#">Riemann</a> and <a href="#">Microsoft System Center Operations Manager</a>.</p>	<p>The data confirms a conventional approach, with Nagios used for monitoring infrastructure and New Relic used for application performance monitoring. Since the survey was conducted, we expect that Datadog and AWS CloudWatch usage has increased.</p>
<p><b>Server infrastructure was monitored by 81 percent of the respondents.</b> Cloud infrastructure, network infrastructure, and applications were monitored less often.</p>	<p>Although DevOps practitioners care about infrastructure, they will not use the same metrics as those valued by traditional IT operations teams.</p>
<p><b>collectd, StatsD and New Relic were the most common products used to collect metrics.</b> That data is being stored most often with <a href="#">Graphite</a>. Although used less often, time series databases <a href="#">InfluxDB</a> and <a href="#">OpenTSDB</a> were also represented in the study.</p>	<p>We expect that container-related agents like Sysdig will show up more often. New monitoring tools like Prometheus have a custom-built time series database at their core.</p>
<p><b>Grafana was most used for visualization</b>, followed by <a href="#">Graphite</a>.</p>	<p>Notably absent from the list is Kibana, which is a commonly used part of the ELK Stack, now called Elastic Stack by its corporate parent.</p>

**TABLE 1:** *New monitoring solutions compare themselves favorably to the commonly used Nagios and Graphite applications.*

we reviewed a [survey](#) James Turnbull conducted in 2015. Although it is a snapshot of people that are already inclined to care about monitoring, it provides many relevant insights.

Expectations of time and effort needed for monitoring is changing. With the current prevalence of automated systems, users often want to reduce the time needed to setup a monitoring tool and trying to find the problem in their stack. While monitoring may always been relatively time consuming, there are approaches that can improve the overall experience.

## What's Different With Containers

To understand how to monitor containers and their related infrastructure, you need to understand what is different about containers. There are aspects of containerized environments that change previously established monitoring practices and the efficiency of traditional monitoring

solutions. Understanding these changes will help explain how vendors are shifting to create new products to address changing metrics and a new, varied team of users involved in monitoring. The monitoring changes that come with containers can be explained in five points:

1. The ephemeral nature of containers.
2. The proliferation of objects, services and metrics to track.
3. Services are the new focal point of monitoring.
4. A more diverse group of monitoring end-users.
5. New mindsets are resulting in new methods.

## Ephemerality and Scale of Containers

Cloud-native architectures have risen to present new challenges. The temporary nature of containers and virtual machine instances presents tracking challenges. As containers operate together to provide microservices, they are in effect a distributed system. While distributed systems are not necessarily transitory at larger scales, they require targeting of many moving parts. This requires new methods of monitoring to make observations about their health. Due to their ephemeral nature and growing scale, it doesn't make sense to track the health of individual containers; instead, you should track clusters of containers and services.

Traditional approaches to monitoring are based on introducing data collectors, agents or remote access hooks into the systems for monitoring. They do not scale out for containers due to the additional complexity they introduce to the thin, application-centric encapsulation of containers. Neither can they catch up to the provisioning and dynamic scaling speed of containers.

In the past, people would look at a server to make sure it was running. They would look at its CPU utilization and allocated memory, and track network bottlenecks with I/O operations. The IT operator would be able to know where the machine was, and easily be able to do one of two things. First, they could point an instrument to that specific location and collect data. In monitoring language, this is called polling a machine. Alternatively, an agent can be installed on the server, which then pushes data to a monitoring tool.

This push approach has achieved popularity because the ephemeral nature of containers and virtual instances makes it difficult to instrument tools to find and poll them. It also reduces the amount of intrusion or tainting of applications. This monitoring approach benefits from the key observability characteristics of containers, and enables solutions that operate efficiently, seamlessly and without intrusion to container execution.

## Proliferation of Objects, Services and Metrics to Track

The explosion of data being generated is a well known phenomenon. Ten years ago, people cared about how to store all that data. More recently, the focus has been on how to best utilize that data without storing it all. With the rise of Internet of Things (IoT) sensors and container adoption, there are now more and more objects than ever to monitor. While there is an instinct to try to corral all these objects into a monitoring system, others are attempting to identify new units of measurement that can be more actionable and easily tracked.

The abundance of data points, metrics and objects that need to be tracked is a serious problem. Streaming data presents many opportunities for real-time analytics, but it still has to be processed and stored. There are technical solutions that can handle the scale, but at significant cost to

both finance and performance. While NoSQL and other next-generation databases have established their place in the IT ecosystem, they are not optimized for this use case; time series databases is a potential solution for storage. However, companies can't just store their log data indefinitely; much of the data is never used. Some older log files are never looked at, motivating users to focus less on log management tools and more on metrics, which is data collected in aggregate or at regular intervals.

Containers present two problems in terms of data proliferation. Compared to traditional stacks, there are more containers per host to monitor and the number of metrics per host has increased. As CoScale CEO Stijn Polfliet describes it, there would traditionally be 150 metrics to track per host: 100 about the operating system and 50 about an application. With containers, you're adding an additional 50 metrics per container and 50 metrics per orchestrator on the host. Considering a scenario where there a cluster is running 100 containers on top of two underlying hosts, there would be over 10,000 metrics to track.

With so much potential data to collect, users focus on metrics. As Honeycomb co-founder and engineer Charity Majors [wrote](#), “Metrics are usually bucketed by rollups over intervals, which sacrifices precious detail

Per Host Metrics Explosion			
Component	# of Metrics for a Traditional Stack	for 10 Container Cluster with 1 Underlying Host	for 100 Container Cluster with 2 Underlying Hosts
<b>Operating System</b>	100	100	200
<b>Orchestrator</b>	n/a	50	50
<b>Container</b>	n/a	500 (50 per container)	5,000 (50 per container)
<b>Application</b>	50	500 (50 per container)	5,000 (50 per container)
<b>Total # of Metrics</b>	<b>150</b>	<b>1,150</b>	<b>10,250</b>

**TABLE 2:** Containers means more metrics than traditional stacks.

about individual events in exchange for cheap storage. Most companies are drowning in metrics, most of which never get looked at again. You cannot track down complex intersectional root causes without context, and metrics lack context.” Even though metrics solve many operations problems, there’s still too many of them, and they’re only useful if they’re actually utilized.

## Services Are the New Focal Point

With a renewed focus on what actually needs to be monitored, there are three areas of focus: the health of container clusters; microservices; and applications.

Assessing clusters of containers — rather than single containers — is a better way for infrastructure managers to understand the impact services will have. While it’s true that application managers can kill and restart individual containers, they are more interested in understanding which clusters are healthy. Having this information means they can deploy the cluster to a different infrastructure or add additional resources to support its optimal operation. Container orchestration solutions help by allowing for efficient scheduling of containers on clusters of hosts.

Many microservices are composed of multiple containers. A common example is a microservice composed of five different containers, each running a different process. If one goes down, another can pop up in its place. However, if this failure is a consistent pattern in the long-term, there will be a degradation of the service. Looking at the microservice as a unit can provide insight into how an entire application is running.

According to Dynatrace Cloud Technology Lead Alois Mayr, in an interview with The New Stack, when looking at application monitoring, “you’re mostly interested in the services running inside containers, rather than the containers themselves. This application-centric information ensures that

the applications being served by the component containers are healthy.” Instead of just looking at CPU utilization, you want to look at CPU time based on a specific application or service; monitor how efficiently services are operating by looking at database queries as well as throughput and response times; and track communication between microservices across containers and hosts.


## More Diverse Group of Monitoring End-Users

The focus on monitoring applications instead of just infrastructure is happening for two reasons. First, a new group of people is involved in the monitoring. Second, applications are more relevant to overall business performance.

Monitoring is still generally reactive, despite progress in recent years. It’s focused on the objectives of the IT team managing the actual infrastructure. This mindset does a disservice to developers because they generally receive data secondhand. Developers are increasingly being held accountable for applications once they have been put into production. As Todd DeCapua and Shane Evans’s [Effective Performance Engineering](#) notes, developers are being asked to “deliver the highest quality and performing product, and provide continuous feedback and optimization recommendations, so other teams can deliver quickly and in fully automated ways.” The DevOps movement has risen, at least in part, as a response to developers’ desire for increased visibility throughout the full application life cycle. Now, DevOps roles are often the full stack managers and operators of applications.

Different roles care about different parts of the monitoring process. Our analysis of the aforementioned Turnbull survey of IT professionals that care about monitoring shows that beyond servers, their areas of interest vary significantly. The data shows a break between the developer and DevOps roles. Based on the survey, 48 percent of developers monitor

# Monitored Environments Differ Per Role



Source: The New Stack Analysis of a 2015 James Turnbull survey. Which of the following best describes your IT job role? What parts of your environment do you monitor? Please select all that apply. Developers, n=94; DevOps, n=278; Sysadmin/Operations/SRE, n=419.

THE NEW STACK

**FIG 1:** *DevOps care more about monitoring cloud infrastructure (65 percent) and application logic (75 percent) as compared to their IT operations-focused peers.*

cloud infrastructure, which is significantly below the 65 percent reported by DevOps roles.

The biggest differences are between DevOps and other IT staff. The data showed that 72 percent of system admins and IT Ops roles monitor networking infrastructure, which is about 20 percentage points higher than the developers and DevOps groups. On the reverse side, 70 percent of developers and 75 percent of DevOps roles monitor application logic, compared to only 59 percent of the IT operations-oriented respondents.

DevOps roles care as much about applications as they do infrastructure, but care more about performance than availability. As James Turnbull writes in [The Art of Monitoring](#):

**“**Orienting your focus toward availability, rather than quality and service, treats IT assets as pure capital and operational expenditure. They aren’t assets that deliver value, they are just assets that need to be managed. Organizations that view IT as a cost center tend to be happy to limit or cut budgets, outsource services, and not invest in new programs because they only see cost and not value.”

Luckily, we’ve seen a trend over the last few years where IT is less of a cost center and more of a revenue center. Increased focus on performance pertains to both IT and the business itself. Regarding IT, utilization of storage or CPU resources is relevant because of their associated costs. From the perspective of the business itself, IT used to only care about availability and mean time to resolve (MTTR). While availability and resolvability are still critical, new customer-facing metrics are also important.

Along with DevOps, the practice of site reliability engineering (SRE) will affect how monitoring tools are used. From this perspective, monitoring will still largely be managed by an operations team, but responsibility for ensuring new applications and services are monitored may be delegated to application developers. Shariq Rizvi, co-founder of Netsil, said in an interview with The New Stack that SREs and DevOps engineers are different from software engineers. He believes SRE teams should split up the management of services, thus creating more specialization. Dan Turchin, co-founder and chief product officer of Neva, said in an interview with The New Stack that he believes DevOps positions are replacing network operations center (NOC) engineers, who were traditionally looking at things from a data center perspective. If the old-school networking stats are being displaced by cloud infrastructure metrics, then this may be true.

The market is responding to this changing landscape. Sysdig added [teams](#) functionality to its software, which lets administrators control access to teams based on their need for dashboards, alerts and data as defined by their orchestration system. While this may solve security issues, the main monitoring benefit is that it simplifies the scope of the data being reviewed, thus letting people focus on the data they need to fulfill their roles. Another example of role-based monitoring is playing out in the Kubernetes world, where the project has been redesigning its dashboard based on the differing needs of application developers, application operators and cluster operators.

## New Mindset, New Methods

Although monitoring is changing to meet the needs of different job roles, it is also moving to a more holistic approach. As Majors [wrote on her blog](#), instead of relying on a fixed set of questions and checks, people should move towards the “observability” of systems. This has to happen because those fixed data points will not provide the needed insights alone. New tools are needed to keep pace and provide the ability to predict what’s going to break. Many of these tools use artificial intelligence techniques.

Observability recognizes that testing won’t always identify the problem. Thus, Majors believes that “instrumentation is just as important as unit tests. Running complex systems means you can’t model the whole thing in your head.” Besides changes in instrumentation, she suggests focusing on making monitoring systems consistently understandable. This means actually defining what the data represents and using the same definitions as your peers do both within and outside the organization. Furthermore, there is a frustration with the need to scroll through multiple, static dashboards. In response, vendors are making more intuitive, interactive dashboards. Companies are even using artificial intelligence to determine what information displays when for each service.

# Approaches to Address the New Reality

Increasing automation and predictive capabilities are common approaches to address new monitoring challenges.

Increasing automation centers around reducing the amount of time it takes to deploy and operate a monitoring solution. According to Steven Acreman, founder and Chief Technical Officer of Dataloop.IO, in an interview with The New Stack, the larger the organization, the more likely it will require customized solutions that can collect and integrate data from all their inputs and applications. Vendors are trying to reduce the number of steps required in the setup process. This might mean that once a monitoring agent is installed on a host, you don't have to think about it. More likely, it means that the tools have the ability to auto-discover new applications or containers.

You also want to automate how you respond to problems. For now, there is a difference between automating certain tasks and automation that takes humans entirely out of the equation. Monitoring systems continue to create automated alerts, but now the alerts are more sophisticated. As James Turnbull says in his [book](#), alerting will be annotated with context and recommendations for escalations. Systems can reduce the amount of unimportant alerts, which mitigates alert fatigue and increases the likelihood that the important alerts will be addressed. For now, the focus is getting the alerts to become even more intelligent. Thus, when someone gets an alert, they are shown the most relevant displays that are most likely to identify the problem. For now, it is simply faster and more effective for a human to review the data and address situations on a case by case basis.

Automating the container deployment process is also related to how you monitor it. It is important to be able to track the setting generated by your

configuration management. This is where container orchestrators can help. Kubernetes, Mesos and Cloud Foundry all enable auto-scaling.

Just as auto-scaling is supposed to save time, so is automating the recognition of patterns. [Big Panda](#), [CoScale](#), [Dynatrace](#), [Elastic Prelert](#), [IBM Bluemix](#), [Netsil](#) and [SignalFx](#) are just a few of the companies that use artificial intelligence to identify patterns and detect anomalies. A common result is that much of the noise created by older monitoring approaches gets suppressed. In an interview with The New Stack, Peter Arjis of CoScale says anomaly detection means you don't have to watch the dashboards as much. The system is supposed to provide early warnings by identifying patterns of behavior among how different services, applications and infrastructure behave.

Dynatrace also uses artificial intelligence to look at downed services or applications to identify the root cause of the problem. Looking at an impacted service, it automatically measures what a good baseline would be to compare that service to, with different microservices having different thresholds. If a threshold is surpassed, then the user of the monitoring system either gets alerted or gets presented with a possible resolution in the dashboard.

## Finding the Most Relevant Metrics

The number of container-related metrics that can be tracked has increased dramatically. Since the systems are more complex and decoupled, there is more to track in order to understand the entire system. This dramatically changes the approach in monitoring and troubleshooting systems.

Traditionally, availability and utilization of hosts is measured for CPUs, memory, I/O and network traffic. Although these are still important for managing IT infrastructure, they do not provide the best frame of reference for evaluating what metrics to collect.

Although there are many different layers in this IT environment, services are a key unit of observation. Service health and performance is directly related to application performance. Services can be defined with common names, with their health and performance benchmarked over time. Services, including microservices running in containers, can be tracked across clusters. Observing clusters of services is similar to looking at the components of an application.

Google's book on [Site Reliability Engineering](#) claims there are four key signals to look at when measuring the health and performance of services: latency, traffic, errors and saturation. Latency describes the time it takes to service requests. Within a container, it can be helpful to look at how slowly API calls are handled. Traffic and errors are both commonly tracked, and refer to the communicating and networking of services and the frequency of errors. Saturation describes how "full" the service is and emphasizes the most constrained resources. It is becoming a more popular way to measure system utilization because service performance degrades as they approach high saturation.

Using this viewpoint, we can see what types of metrics are most important throughout the IT environment. Information about containers is not an end unto itself. Instead, container activity is relevant to tracking infrastructure utilization as well as the performance of applications and infrastructure. Metrics about the saturation and latency of requests within a container are most relevant. Metrics about the health of individual containers will continue to be relevant. However, in terms of managing containers, measuring the health of clusters of containers will become more important.

It's important to remember that you're not just monitoring containers, but also the hosts they run on. Utilization levels for the host CPU and memory can help optimize resources.

## Questions to Ask When Deciding What Metrics to Monitor

	Questions	Sample Metrics
<b>Microservice</b> In general, there is one process to track per container.	Where are new services deployed? What percentage of time is the service reachable? How many requests are enqueued?	Average percentage of time a request-servicing thread is busy. Number of enqueued requests. Percentage of time a service is reachable.
<b>Application</b> Multiple microservices running simultaneously constitute an application.	Do the databases respond quickly? Are the message queues fast enough? How does heap memory usage change over time? Are my application services responsive?	Query execution frequency, response time and failure rate. Response time, failure rate.
<b>Container</b> Separate from the underlying process being run within it, containers are also monitored.	How responsive are the processes within the container? Which images have been deployed? Are specific containers associated with over-utilization of hosts?	CPU throttle time. Container disk I/O. Memory usage. Network (volume, dropped packets).
<b>Container Cluster</b> Multiple containers deployed to run as a group. Many of the metrics for individual containers can also be summarized.	Are your clusters healthy and properly sized? Can applications be effectively run using fewer nodes?	Percentage of clusters remaining operational compared to those originally deployed.
<b>Host</b> Also called a node, multiple hosts can support a cluster of containers.	Do changes in utilization indicate a problem with a process or application?	Percentage of total memory capacity in use. Percentage of time CPUs are utilized.
<b>Infrastructure</b> Broadly speaking, this is the cloud in which the hosts are running.	How much does it cost to run each service or deployment? What is the ratio of microservices and/or containers per instance?	Network traffic. Utilization of databases, storage, and other shared services.
<b>End User</b> The end goal of the entire system is to serve this group.	What is the average web response time experienced by users per region?	Response time. Number and percentage of user actions that failed.

**TABLE 3:** Saturation and latency related metrics are the most relevant when monitoring microservices-based applications. Instead of looking at individual services and containers, dashboards and alerts should focus on their operation in aggregate.

As Sematext DevOps Evangelist Stephan Thies [wrote](#), “when the resource usage is optimized, a high CPU utilization might actually be expected and even desired, and alerts might make sense only for when CPU utilization drops (service outages) or increases for a longer period over some max limit (e.g., 85%).”

In the past, it was possible to benchmark host performance based on the number of applications running on it. If environments weren’t dynamic, with virtual instances being spun up and down, then it would be possible to count the number of containers running and compare it to historical performance. Alas, in dynamic environments, cluster managers are automatically scheduling workloads, so this approach is not possible. Instead, observing the larger IT environment for anomalies is becoming a way to detect problems.

## The Next Steps

The biggest changes in IT monitoring are the new groups involved and the new metrics they are using. IT operations still care about availability and cost optimization. DevOps and application developers focus on the performance of services. Everyone, especially the chief information officer, cares about the impact on business operations and customer interactions.

Of course, there are new metrics that have to be monitored. The [Identifying and Collecting Container Data](#) chapter provides an overview of how to collect this data. All of these metrics can be collected in different ways. [Classes of Container Monitoring](#) details the different components of an effective monitoring stack. From collection to logging to visualization, there are unique technical challenges to monitoring containers and microservices. Looking at next steps, [The Right Tool for the Job: Picking a Monitoring Solution](#) provides important criteria to think about.

# CREATING ANALYTICS-DRIVEN SOLUTIONS FOR OPERATIONAL VISIBILITY




In this discussion with Canturk Isci, master inventor and researcher at IBM, we talk about how monitoring practices can be elevated to be more than just responsive systems. By using principles of operational visibility to create analytics-driven products, monitoring can go beyond basic reporting and begin to solve real-world problems. The analytics-driven component contributes to a monitoring system that can be proactive by using analytical models for how containers should behave. We later discuss some of the differences between monitoring traditional and even VM-based environments, and newer container-based environments. Isci talks about a growing base of products that package everything for you, including the components necessary for visibility, monitoring and reporting. IBM has a wide variety of monitoring capabilities, many of which are focused on providing monitoring to IBM Bluemix and container services. [Listen on SoundCloud](#) or [Listen on YouTube](#)



**Canturk Isci** is a researcher and master inventor in the IBM T.J. Watson Research Center, where he also leads the Cloud Monitoring, Operational and DevOps Analytics team. He currently works on introspection-based monitoring techniques to provide deep and seamless operational visibility into cloud instances, and uses this deep visibility to develop novel operational and DevOps analytics for cloud. His research interests include operational visibility, analytics and security in cloud, DevOps, configuration analytics, energy-efficient and adaptive computing.

# CLASSES OF CONTAINER MONITORING

by **BRIAN BRAZIL**



Before we talk about container monitoring, we need to talk about the word “monitoring.” There are a wide array of practices considered to be monitoring between users, developers and sysadmins in different industries. Monitoring — in an operational, container and cloud-based context — has four main use cases:

- Knowing when something is wrong.
- Having the information to debug a problem.
- Trending and reporting.
- Plumbing.

Let’s look at each of these use cases and how each obstacle is best approached.

## **Knowing When Something is Wrong**

Alerting is one of the most critical parts of monitoring, but an important question to ask is: What problem is worth waking up an engineer in the middle of the night to look at? It’s tempting to create alerts for anything

that's a red flag or even slightly troublesome, but this can quickly lead to alert fatigue.

Let's say you're running a set of user-facing microservices, and you care about the latency of requests. Would central processing unit (CPU) usage on each machine be useful to alert on? The alert will likely flag you that you're running out of CPU capacity on the machine. It will also have false positives when background processes take a little longer than usual, and false negatives for deadlocks or not having enough threads to use all CPUs.


The CPU is the potential cause of the problem, and high latency is the symptom you are trying to detect. In [My Philosophy on Alerting](#), Rob Ewaschuk points out that there are many potential causes, and it's difficult to enumerate all of them. It's better to alert on the symptoms instead, as it results in fewer pages that are more likely to present a real problem worth waking someone up over. In a dynamic container environment where machines are merely a computing substrate, alerting on symptoms rather than causes goes from being a good idea to being essential.

## Having the Information to Debug a Problem

Your monitoring system now alerts you to the fact that latency is high. Now what do you do? You could go login to each of your machines, run the [top command](#), check syslog and start tailing application logs. That's not going to get you anywhere fast though, and it will lose effectiveness as your traffic and infrastructure grows. What your monitoring needs to provide is a way for you to approach problems methodically, giving you the tools you need to narrow down issues.

Microservices can typically be viewed as a tree, with remote procedure calls (RPCs) flowing from the top to the bottom. A problem of high latency in a service is usually caused by a delay in that service or one of its backends. Rather than trying to get inspiration from hundreds of graphs

# Microservice Routing



Source: Brian Brazil

THE NEW STACK

**FIG 1:** Component routing in a microservice.

on a dashboard, you can go to the dashboard for the root service and check for signs of overload and delay in its backends. If the delay is in a backend, you repeat the process until you find the service responsible.

That process can be taken a step further. Just like how your microservices compose a tree, the subsystems, libraries and middleware inside a single microservice can also be expressed as a tree. The same symptom identification technique can then be applied to further narrow down the issue. To continue debugging from here, you'll likely use a variety of tools to dig into the process internals, investigate patterns in request logs and cross-correlate requests across machines.

## Trending and Reporting

Alerting and debugging tend to be on the timescale of minutes to days. Trending and reporting care about the weeks-to-years timeframe.

A well-used monitoring system collects all sorts of information, from raw hardware utilization and counts of API requests to high-level business metrics. There are the obvious use cases, such as provisioning and capacity planning to be able to meet future demand, but beyond that there's a wide selection of ways that data can help make engineering and business decisions.

Knowing how similar requests are to each other might point to the benefit of a cache, or it might help argue for removing a cache for simplicity. Knowing how each request uses your limited resources can help determine your pricing model. Cross-service and cross-machine statistics can help you spend your time on the best potential optimizations. Your monitoring systems should empower you to make these analyses possible.

## Plumbing

When you have a hammer, everything starts to look like a nail.

Plumbing is different from the other use cases, as it's about getting data from system A to system B, rather than directly supporting responsive decision making. An example might be sending data on the number of sales made per hour to a business intelligence dashboard. Plumbing is about facilitating that pipeline, rather than what actions are taken from the end result. It's not necessarily monitoring; however, it's often convenient to use your monitoring system to move some data around to where it needs to go.

If building a tailored solution from scratch could take weeks, and it's effectively free to use your monitoring system for the same thing, then why not? When evaluating a monitoring system, don't just look at its ability to do graphing and alerting, but also how easy it is to add custom data sources and extract your captured data later.

# Classes of Monitoring

Now that we've established some of what monitoring is about, let's talk about the data being inserted into our monitoring systems. At their core, most monitoring systems work with the same data: events. Events are all activities that happen between observation points. An event could be an instruction being executed, a function call being made, a request being routed, a remote call procedure (RPC) being received or a response being returned. Events have contextual information, such as what triggered them and what data they're working with.


We're going to look at four different ways to use events; each approach makes different tradeoffs and gives you a different view of the system. A complete monitoring system will have aspects of each approach.

## Metrics

Metrics, sometimes called time series, are concerned with events aggregated across time. They count how often each type of event happens, how long each type of event takes and how much data was processed by the event type.

Metrics largely don't care about the context of the event. You can add context, such as breaking out latency by HTTP endpoint, but then you need to spend resources on a metric for each endpoint. In this case, the number of endpoints would need to be relatively small. This limits the ability to analyze individual occurrences of events; however, in exchange, it allows for tens of thousands of event types to be tracked inside a single service. This means that you can gain insight into how code is performing throughout your application. We're going to dig a bit deeper into the constituent parts of metrics-based monitoring. If you're only used to one or two systems, you may not be aware of the possibilities and tradeoffs that can be made.

# Monitoring Metrics Pipeline



Source: Brian Brazil

THE NEW STACK

**FIG 2:** The architecture of gathering, storing and visualizing metrics.

## Collection

Collection is the process of converting the system state and events into metrics, which can later be gathered by the monitoring system. Collection can happen in several ways:

1. **Completely inside one process.** The [Prometheus](#) and [Dropwizard](#) instrumentation libraries are examples; they keep all state in memory of the process.
2. **By converting data from another process into a usable format.** [collectd](#) and [Agentless System Crawler](#) do this by pulling data from the [proc filesystem](#).
3. **By two processes working in concert: one to capture the events and the other to convert them into metrics.** [StatsD](#) is an example, where each event is sent from an application over the network to StatsD.

## Ingestion

Ingestion takes metrics from collection and feeds them into the monitoring system. This can be a multi-stage process involving a queueing system, such as [Apache Kafka](#), or a simple data transfer directly from collection. It's at this point that the push versus pull debate must be mentioned. Both

approaches have advantages and disadvantages. We can't cover the extent of this debate in these pages, but the short version is that both approaches can be scaled and both can work in a containerized environment.

## Storage

Once data is ingested, it's usually stored. It may be short-term storage of only the latest results, but it could be any amount of minutes, hours or days worth of data storage.

Once stored data goes beyond what easily fits in memory on one machine, there's operational and reliability tradeoffs to be made, and again there are pros and cons based on what the organization requires from their monitoring data. Persisting data beyond the lifetime of a process on disk implies either a need for backups or a willingness to lose data on machine failure.

Spreading the data among multiple machines brings with it the fundamental challenges of distributed systems. It's not difficult to end up with a system where existing data is safe, but new data cannot be ingested and processed.

## Processing and Alerting

Data isn't of much use if you don't do anything with it. Most metrics systems offer some way to do math on ingested data, and usually also offer a way to alert humans of anomalous conditions. This may happen as the data is ingested or as a separate asynchronous process.

The sophistication of processing between solutions varies greatly. On one end, [Graphite](#) has no native processing or alerting capability without third-party tools; however, there's basic aggregation and arithmetic possible when graphing. On the other end, there are solutions like Prometheus or Sysdig with not only a fully-fledged processing and alerting systems, but also an additional aggregation and deduplication system for alerts.

## Visualization

Alerts arriving at your pager is fine, but for debugging, reporting and analysis you want dashboards to visualize that data.

Visualization tools tend to fall into three categories. At the low end, you have built-in ways to produce ad-hoc graphs in the monitoring system itself. In the middle, you have built-in dashboards with limited or no customization. This is common with systems designed for monitoring only one class of system, and where someone else has chosen the dashboards you're allowed to have. Finally, there's fully customizable dashboards where you can create almost anything you like.


## How They Fit Together

Now that you have an idea of the components involved in a metrics monitoring system, let's look at some concrete examples of the tradeoffs made by each.

## Nagios

The [Nagios](#) server usually calls out to scripts on hosts — called checks — and records if they work according to their exit code. If a check is failing, it sends out an alert. Visualization is typically offered by a separate built-in dashboard. It can ingest 1KB of data, including metrics (called “perfdata”), from the script and pass it on to another monitoring system.

**FIG 3:** Metrics handling with Nagios.



Nagios is designed for static setups, which requires a restart to load a new configuration. Its limited processing, focus on host-based monitoring, and ability to only handle small amounts of metrics data makes it unsuitable for monitoring in a container environment. However, it remains useful for basic blackbox monitoring.


## collectd, Graphite and Grafana

Many common monitoring stacks combine several components together. A [collectd](#), [Graphite](#) and [Grafana](#) combination is an example of such. collectd is the collector, pulling data from the kernel and third-party applications such as MySQL. To collect custom metrics from your own applications, you'd use the StatsD protocol, which sends user data protocol (UDP) packets to collectd for individual events. collectd sends metrics to [Carbon](#), which uses a [Whisper](#) database for storage. Finally, both Graphite and Grafana themselves can be used for visualization.

The StatsD approach to collection is limiting in terms of scale; it's not unusual to choose to drop some events in order to gain performance. The collectd per-machine approach is also limiting in a containerized environment. For example, if there are MySQL containers dynamically deployed, then the per-machine collectd needs its configuration updated each time.

**FIG 4:** An example monitoring stack composed of [collectd](#), [Graphite](#), [Grafana](#).

## Example Monitoring Architecture




As alerting is not included, one approach is to have a Nagios check for each individual alert you want. The storage for Graphite can also be challenging to scale, which means your alerting is dependent on your storage being up.

## Prometheus

Prometheus takes a different approach than our previous examples. Collection happens where possible inside the application. For third-party applications where that's not possible, rather than having one collector per machine, there's one exporter per application. This approach can be easier to manage, at the cost of increased resource usage. In containerized environments like Kubernetes, the exporter would be managed as a sidecar container of the main container. The Prometheus server handles ingestion, processing, alerting and storage. However, to avoid tying a distributed system into critical monitoring, the local Prometheus storage

**FIG 5:** Metrics handling in Prometheus.

## Prometheus Architecture



is more like a cache. A separate, non-critical distributed storage system handles longer term storage. This approach offers both monitoring reliability and durability for long-term data.


While Prometheus decides what alerts to fire, it does not send emails or pages to users. Alerts are, instead, sent to an Alertmanager, which deduplicates and aggregates alerts from multiple Prometheus servers, and sends notifications.

## Sysdig Cloud

The previous sections show how various open source solutions are architected. For comparison, this section describes the architecture of Sysdig Cloud, a commercial solution. Starting with instrumentation, Sysdig Cloud uses a per-host, kernel level collection model. This instrumentation captures application, container, statsd and host metrics with a single collection point. It collects event logs such as Kubernetes scaling, Docker container events and code pushes to correlate with metrics. Per-host

**FIG 6:** Taking a look at Sysdig's commercial architecture.

## Sysdig Cloud Architecture



agents can reduce resource consumption of monitoring agents and require no modification to application code. It does, however, require a privileged container.

The Sysdig Cloud storage backend consists of horizontally scalable clusters of Cassandra (metrics), MySQL (events), and Redis (intra-service brokering). Building on these components gives high reliability and scale to store years of data for long term trending and analysis. All data is accessible by a REST API. This entire backend can be used via Sysdig's cloud service or deployed as software in a private cloud for greater security and isolation. This design allows you to avoid running one system for real-time monitoring and another system for long term analysis or data retention.

In addition to handling metrics data, Sysdig Cloud also collects other types of data, including events logs from Kubernetes and Docker containers, and metadata from orchestrators such as Kubernetes. This is used to enrich the information provided from metrics, and it's not unusual for a metrics system to have integrations beyond what's purely required for metric processing.

## Logs

Logs, sometimes called event logs, are all about the context of individual events. How many requests went to an endpoint? Which users are using or calling an endpoint?

Logs make the opposite tradeoff to metrics. They don't do any aggregation over time. This limits them to tracking around fifty to a hundred pieces of information per event before bandwidth and storage costs tend to become an issue. Even with this limitation, logs usually allow you to find patterns in individual requests, such as if particular users are hitting expensive code paths.

It's important to distinguish the type of logs you are working with, as they have a variety of different uses and reliability requirements:

- **Business and transaction logs:** These are logs you must keep safe at all costs. Anything involved with billing is a good example of a business or transaction log.
- **Request logs:** These are logs of every request that comes through your system. They're often used in other parts of the system for optimization and other processing. It's bad to lose some, but not the end of the world.
- **Application logs:** These are logs from the application regarding general system state. For example, they'll indicate when garbage collection or some other background task is completed. Typically, you'd want only a few of these log messages per minute, as the idea is that a human will directly read the logs. They're usually only needed when debugging.
- **Debug logs:** These are very detailed logs to be used for debugging. As these are expensive and only needed in specialized circumstances, they have lower reliability requirements than application logs.

The next time someone talks to you about logs, think about which type of logs they're talking about in order to properly frame the conversation.

## Profiling

Profiling has the same advantages of metrics and logs. It lets you see data about individual events throughout the entire application. The disadvantage is that this tends to be very expensive to do, so it can only be applied tactically.

For example, logs have told you that a user is hitting an expensive code path, and metrics have let you narrow down which subsystem is the likely

culprit. Your next step is to profile that subsystem and see in which exact lines of code the CPU is being spent.

There are a variety of Linux [profiling tools](#), including [eBPF](#), [gdb](#), [iostop](#), [strace](#), [tcpdump](#) and [top](#). There are also commercial options, like [Sysdig](#), which combine functionality of several of these tools into one package. You can use some of these tools on an ongoing basis, in which case it would fall under metric or logs.

## Distributed Tracing


Let's say you have a system with a frontend running at 310ms latency in the 95th percentile. You receive an alert saying the frontend 95th percentile latency has increased to 510ms! What do you think you'll see in the 95th percentile latency of the culprit backend?

The answer is that you might see an increase of the same size as on the frontend, but you might not. There could be no change, or even a decrease in latency. It all depends on the correlations of the latencies. Remember, the 95th percentile is effectively throwing 95 percent of data away, so you won't notice changes outside of that 5 percent.

What's going on here isn't obvious from the latency graphs, and that's where distributed tracing comes in. It's a form of logging and profiling. It is particularly useful in environments such as those using containers and microservices with a lot of inter-service communication.

How it works is that each individual incoming request gets a unique identifier. As the request passes through different services on different machines, that information is logged with the identifier. Each request is then stitched back together from the logs to see exactly where time was spent for each request. Due to cost of the logging, it's sometimes only possible to trace a subset of incoming requests.

## System Latency (at 95th Percentile) Mapping



Source: Brian Brazil

THE NEW STACK

**FIG 7:** An example system's latency mapping.

The result is a visualization of when each backend in your tree of services was called, allowing you to see where time is spent, what order requests are made in and which RPCs are on the critical path. Applied to the example in Figure 6, you'd notice that all the fast requests only hit Backend 1, while the slow requests are hitting both backends. This would tip you off that it's the logic about communicating with Backend 2 that you need evaluate. While this is a simple example, imagine how long it'd take you to figure out if you had dozens of services without distributed tracing.

## Conclusion

In this article, we've covered the use cases for monitoring, which should help you understand the problems that can be solved with monitoring. We learned about the four different ways for using events: metrics, logs, profiling and distributed tracing. In breaking down the metrics-based approach, we looked at how data is collected, ingested, stored, processed, alerted and visualized.

Now that you have a better feel for the types of monitoring systems and the problems they solve, you will be able to thoroughly evaluate many

different solutions. There are many approaches to designing a monitoring system, and each have their own advantages and disadvantages. When looking to evaluate a monitoring solution, first assess whether it's primarily based on metrics, logs, profiling or distributed tracing. From there, see what features it has that'll fit into your overall monitoring strategy, in terms of alerts requiring intelligent human action, the information you need to debug, and to integrate with your systems.

Each solution has its pros and cons, and you'll almost certainly need more than one tool to create a comprehensive solution for monitoring containers.

# BUILDING ON DOCKER'S NATIVE MONITORING FUNCTIONALITY




In this podcast with John Willis, director of Ecosystem Development at Docker, we discuss how monitoring applications and systems changes when you introduce containers. Willis talks about functionally treating containers as having multiple perimeters, which involves treating the host and container as isolated entities. Many of the basic functions of monitoring remain the same, but there's way more metadata involved with containers. That means taking advantage of features such as Docker's APIs, and building off of that core functionality. The Ecosystem Technology Partner (ETP) program continues to produce products and services that expand Docker monitoring functionality through an ecosystem of new solutions. Willis also talks about the recently released Docker InfraKit, a toolkit for self-healing infrastructure, and how that will impact container monitoring and management. [Listen on SoundCloud](#) or [Listen on YouTube](#)



**John Willis** is the director of Ecosystem Development for Docker, which he joined after the company he co-founded (*SocketPlane*, which focused on SDN for containers) was acquired by Docker in March 2015. Previous to founding *SocketPlane* in Fall 2014, John was the chief DevOps evangelist at Dell, which he joined following the *Enstratius* acquisition in May 2013. He has also held past executive roles at Chef and Canonical. John is the author of seven IBM Redbooks and is co-author of [The Devops Handbook](#), along with authors Gene Kim and Jez Humble.

# IDENTIFYING AND COLLECTING CONTAINER DATA

by **RUSS MCKENDRICK**



With the introduction of containers and microservices, monitoring solutions have to handle more ephemeral services and server instances than ever before. And while the infrastructure landscape has changed, operations teams still need to monitor the same information on the central processing unit (CPU), random access memory (RAM), hard disk drive (HDD), network utilization, and the availability of application endpoints.

While you can use an older tool or existing monitoring service for traditional infrastructures, there are newer cloud-based offerings that can ensure monitoring solutions are as scalable as the services being built and monitored. Many of these cloud-based and self-hosted tools are purpose-built for containers. No matter what solution or service you use, you still need to know how you're going to collect the metrics you're looking to monitor.

There are quite a few options when it comes to collecting metrics from your containers. This article overviews some of the the services available for container monitoring. We've included a mixture of self-hosted open

source solutions and commercial cloud-based services to reflect the current landscape.

However, it's important to keep in mind that we've included a relatively small amount of examples compared to the wide variety of solutions available. These examples are meant to be illustrative of a few different ways to approach the problem of collecting metrics.

## Docker

The Docker engine itself provides access to most of the core metrics we're looking to collect as part of its native monitoring functions. Running the `docker stats` command will give you access to CPU, memory, network

CONTAINER	CPU %	MEM USAGE / LIMIT	MEM %	NET I/O	BLOCK I/O	PIDS
5849984035e5	0.97%	14.45 MiB / 1.953 GiB	0.72%	1.006 kB / 648 B	0 B / 0 B	5
ce23b2496b5c	0.12%	14.27 MiB / 1.953 GiB	0.71%	1.654 kB / 648 B	0 B / 0 B	5
ca8961914c98	0.05%	14.27 MiB / 1.953 GiB	0.71%	2.232 kB / 648 B	0 B / 0 B	5
743af45ff75c	0.08%	14.22 MiB / 1.953 GiB	0.71%	2.74 kB / 648 B	0 B / 0 B	5
07bdeeeee2ba	1.16%	14.31 MiB / 1.953 GiB	0.72%	107.6 kB / 131.3 kB	0 B / 0 B	5

**FIG 1:** Running the `docker stats` command.

and disk utilization for all of the containers running on your host.

The data streams automatically and is useful if you need to get a quick overview of your containers at any given moment. There are a few flags you can add, for example:

- The flag `--all` shows you stopped containers, though you won't see any metrics for them.
- The flag `--no-stream` displays the output of the first run and then stops the stream of metrics.

There are some disadvantages. For one, the data isn't stored anywhere — you can't go back and review the metrics. It can also be tedious watching a constantly refreshing terminal with no points of reference. It's

hard to spot what's going on in the data. Luckily, the `docker stats` command is actually an interface to the stats API endpoint. The stats API exposes all of the information in the stats command and more. To view for yourself, run the following command:

```
curl --unix-socket /var/run/docker.sock...
http://containers/container_name/stats
```

As you'll be able to see from the output, there is a lot more information being returned, all wrapped in a JavaScript Object Notation (JSON) array, which is ready to be ingested by third-party tools.

## cAdvisor


[cAdvisor](#) is a container monitoring tool from Google with native support for Docker containers. It is a daemon that collects, aggregates, processes, and exports information about running containers.

The best way to think of cAdvisor is that it's a graphical version of what you get if you were to run the `docker stats --all` command.

```
docker run \
--volume=/:/rootfs:ro \
--volume=/var/run:/var/run:rw \
--volume=/sys:/sys:ro \
--volume=/var/lib/docker/:/var/lib/docker:ro \
--publish=8080:8080 \
--detach=true \
--name=cadvisor \
google/cadvisor:latest
```

cAdvisor is easy to get up and running, as it is shipped in a single container. All you have to do is run the above command to launch a cAdvisor container and expose the web interface on port 8080.

Once launched, cAdvisor will hook itself into the Docker daemon running on your host machine and start gathering metrics for all of your running containers, including the cAdvisor container itself. Opening <http://localhost:8080/> in a browser will take you straight to the web interface.



**FIG 2:** Accessing the *localhost* web interface.

As you can see from Figure 2, there is a minute's worth of data being streamed in real time; however, you can't view any metrics further back with just a standard installation of cAdvisor. Luckily, Google has you covered by including options to export the data from cAdvisor into time series databases such as [Elasticsearch](#), [InfluxDB](#), [BigQuery](#) and [Prometheus](#).

To sum up, cAdvisor is a great tool for getting quick insight into what's happening with your running containers. It is simple to install and gives you more granular metrics than Docker does out of the box. It can also act as a monitoring agent for other tools in the same way that Zabbix or Nagios agents were used in the past.


## Prometheus

[Prometheus](#) is an open source monitoring system and time series database originally built by SoundCloud. It is currently hosted by the [Cloud Native Computing Foundation](#) (CNCF), alongside other projects such as [Kubernetes](#) and [OpenTracing](#). Upon reading the description of the service, it may sound like quite a traditional server and agent setup; however, Prometheus works differently. Rather than having its own agent, it scrapes data from the data points on hosts, storing it in its own time series database.

At the end of October 2016, a pull request to [expose a metrics endpoint](#) in the Docker Engine itself was merged. This likely means that an upcoming release of Docker will have support for Prometheus to scrape your container metrics out of the box; however, until that release, you need to use an intermediary service such as cAdvisor. As of 1.13, the Docker engine supports an optional “/metrics” Prometheus endpoint. It’s important to note that this currently exposes internal Docker Engine metrics versus container metrics. There is [ongoing discussion](#) about expanding this to cover container metrics, and possibly replacing the Stats API altogether.

Prometheus’s biggest strength is as a data source. You can front the data being scraped by Prometheus with [Grafana](#), which has supported Prometheus since mid-2015, and is now the recommended frontend for Prometheus. Like most of the tools we are looking at, Grafana can also be launched as a container.

Once up and running, the only configuration needed is to add your Prometheus URL as a data source, then import one of the predefined [Prometheus dashboards](#).



**FIG 3:** Metrics stored in Prometheus and displayed with Grafana.

The dashboard in Figure 3 displays over an hour of metrics from cAdvisor, which are stored in Prometheus and rendered in Grafana. As Prometheus

is essentially taking snapshots of the current cAdvisor state, it's also recording metrics for containers long since removed.

Prometheus also has alerting capabilities. By using the inbuilt reporting language, you can create alerts such as the following:

```
ALERT InstanceDown
  IF up == 0
    FOR 5m
    LABELS { severity = "page" }
    ANNOTATIONS {
      summary = "Instance {{ $labels.instance }} down",
      description = "{{ $labels.instance }} of job {{ $labels.job }} has been down for more than 5 minutes.",
    }
```

Once your alert has been written and deployed on the Prometheus server, you can use the Prometheus [Alertmanager](#) to route your alerts. In the example above, we have assigned a label of **severity = "page"**. Alertmanager will intercept the alert and forward it to a service such as [PagerDuty](#), [OpsGenie](#), a [Slack](#) or [HipChat](#) channel, or any number of [different endpoints](#).

Prometheus is a powerful platform and works brilliantly as a middleman between different technologies. It is easy to start off with a basic installation, like the one described above, then expand it, giving a single pane of glass view of both your container and host instances.

## Agentless System Crawler

[Agentless System Crawler](#) (ASC) is a cloud monitoring tool from IBM with support for containers. It collects monitoring information from running


containers including metrics, system state and configuration information. ASC provides deep visibility into containers, beyond just utilization and performance metrics, for security and configuration analytics. It is designed as a pipeline for building collection plugins for container features, function plugins for on-the-fly data aggregation or analysis, and output plugins for target monitoring and analytics endpoints. Provided plugins include traditional data collection features such as utilization metrics, processes running in the container and listening on a port, and data about the configuration file and application packages.

ASC can be deployed either as a Python package or as a single container with privileges. For example, ASC operates by interacting with the Docker daemon, container cgroups and namespaces. Here is how you can launch ASC as a container:

```
docker run \
    --privileged \
    --net=host \
    --pid=host \
    -v /cgroup:/cgroup:ro \
        -v /var/lib/docker:/var/lib/docker:ro \
    -v /sys/fs/cgroup:/sys/fs/cgroup:ro \
    -v /var/run/docker.sock:/var/run/docker.sock \
    -it crawler --crawlmode OUTCONTAINER ${CRAWLER_ARGS}
```

Once running, it subscribes to the Docker daemon for events and periodically monitors all the containers running in the system. ASC can be easily extended with plugins for custom data collection, which helps to collect both container and environment information. Each plugin simply exports a **crawl()** method for their specific monitored feature; ASC iterates over each **crawl()** function implemented by the collection

plugins enabled in its deployed configuration. ASC can be currently configured to emit data to Kafka, Graphite, filesystem and the console as endpoints, and allows for additional output plugins. Figure 4 shows an example view from an ASC configuration, emitting container data to a Graphite/Grafana endpoint.



**FIG 4:** View of ASC Container Monitoring with Graphite/Grafana Endpoint.

One use case of ASC monitoring is to gain deeper insight into container execution. The following example shows the basic output frame of ASC when configured to collect the process and connection features of a container in addition to base metrics. Based on this frame, we can see that this is a container running at 100 percent CPU utilization, with a single Python process listening on port 5000.

```
metadata    "metadata"      {"id":"ef0fc187dc0a4b85a3cb", ...}
cpu        "cpu-0"          {... , "cpu_util":100.0}
process    "python/1"        {"cmd":"python web.py",
                           "python","pid":74, ...}
connection "74/127.0.0.1/5000" {"pid":74,...,"connstatus":"LISTEN"}
```

This enables users to couple monitoring with system state for better understanding of container behavior and to drive additional analytics. For

instance, we can dig deeper into the Python process used in the container and track the underlying Python packages used by the application by enabling another feature plugin, “python-package” in this case, which gives us the additional insight on the used packages and versions in the application:

```
python-package "click"           {"ver":"6.6","pkgname":"click"}  
python-package "Flask"            {"ver":"0.11.1","pkgname":"Flask"}  
python-package "itsdangerous"    {"ver":"0.24","pkgname":"itsdangerous"}  
python-package "Jinja2"           {"ver":"2.8","pkgname":"Jinja2"}  
python-package "pip"              {"ver":"8.1.2","pkgname":"pip"}  
python-package "setuptools"       {"ver":"20.8.0","pkgname":"setuptools"}  
python-package "Werkzeug"         {"ver":"0.11.11","pkgname":"Werkzeug"}
```

In summary, ASC is a simple tool to set up and use; it’s similar to cAdvisor for gaining visibility into running containers. In addition to basic monitoring metrics, ASC puts a lot of emphasis on deep visibility into container state and configuration to enable monitoring, security and other analytics solutions. It provides a simple, plugin-based approach for extensibility of data collection and support for different endpoints.

## Sysdig

[Sysdig](#) comes in two different versions. The first is an open source version which installs a kernel module on your host machine. The second is a cloud and on-premises solution called [Sysdig Cloud](#), which uses the open source version and streams the metrics it collects to Sysdig’s own servers.

## Open Source Sysdig

The open source version of Sysdig is like running the `docker stats` command on steroids. The service hooks into the host’s kernel, which

means it doesn't entirely rely on getting metrics from the Docker daemon.

Using Csysdig, the built-in [ncurses](#)-based command interface, you can view all sorts of information on your host. For example, running the command **csysdig -vcontainers** gives you the view in Figure 5.

The screenshot shows a terminal window titled "root@monitoring: ~ — ssh root@104.236.12.84 — 124x16". The title bar also includes "Viewing: Containers For: whole machine", "Source: Live System", and "Filter: container.name != host". The main area displays a table of running containers with the following columns: CPU, PROCS, THREADS, VIRT, RES, FILE, NET, ENGINE, IMAGE, ID, and NAME. The table lists eight containers, each running a "docker" engine with the "russmckendrick/cluster" image. The last row shows a summary: "F1Help F2Views F4Filter F5Echo F6Dig F7Legend F8Actions F9Sort F12Spectro CTRL+FSearch p Pause". The status bar at the bottom right shows "1/7(14.3%)".

CPU	PROCS	THREADS	VIRT	RES	FILE	NET	ENGINE	IMAGE	ID	NAME
0.50	3	3	186M	31M	78K	0.00	docker	russmckendrick/cluster	a7d6911b8125	amazing_brattai
0.50	3	3	199M	34M	78K	0.00	docker	russmckendrick/cluster	178275f36280	furious_cray
0.00	3	3	120M	22M	462	0.00	docker	russmckendrick/cluster	dbe3a1923a52	angry_euler
0.00	3	3	120M	22M	710	0.00	docker	russmckendrick/cluster	b7f6d1bf9321	cranky_austin
0.00	3	3	120M	22M	710	0.00	docker	russmckendrick/cluster	2286f1843d80	sad_galileo
0.00	3	3	120M	22M	710	0.00	docker	russmckendrick/cluster	9108351af7db	zen_brown
0.00	3	3	120M	22M	710	0.00	docker	russmckendrick/cluster	eb9a5a8b54fa	loving_lamport

**FIG 5:** Using Sysdig's command interface.


As you can see, this shows all containers running on the host. You can also drill down into a container to see what resources individual processes are consuming. Much like running the **docker stats** command and using cAdvisor, the open source version of Sysdig is about getting a real-time view of your containers; however, you are able to record and replay system activity using the following:

- The command **sysdig -w trace.scap** records the system activity to a trace file.
- The command **csysdig -r trace.scap** replays the trace file.

The open source version of Sysdig is not your traditional monitoring tool; it allows you to drill down into your containers, giving access to a wider range of information. It also allows you to add orchestration context by hooking directly into your orchestrator, thereby allowing you to troubleshoot by pod, cluster, namespace, etc.

## Sysdig Cloud

Sysdig Cloud takes all of the data captured by the open source version and presents it with dashboards and that have alerting functionality. In Figure 6, you can see a dashboard showing a real-time view of container utilization. You're also able to drill down into individual processes.



**FIG 6:** View of dashboards in Sysdig Cloud.

For example, running Wordpress using Docker Compose launches two containers: one running Apache with PHP for the main Wordpress code, and the other running MySQL for the database. Using Sysdig Cloud, you can apply a filter which gives insight into how your database is performing, allowing you to spot such things as slow queries, or which query is being run the most.

There are similar filters for most common services, such as [NGINX](#), [PHP-FPM](#), [PostgreSQL](#), [MongoDB](#) and [Apache Tomcat](#). These filters and metric alerts can be added at any time. These alerts can be delivered to services such as Slack or PagerDuty when triggered, and a system activity

recording will automatically generate in real time. Having a snapshot with that level of detail is invaluable when looking into incidents.

## Other Tools for Collecting Data

As mentioned before, there are quite a few services that can collect, collate and report on metrics from your containers and host instances. We've created an overview of other tools available; there's many more listed in the Container Monitoring Directory.

Other Tools for Collecting Data	
Vendor	Solution Description
CoScale	<b>CoScale</b> is a full-stack monitoring cloud service that monitors the responsiveness of applications. It combines server and container resource metrics with application-specific performance metrics. Its lightweight agents have little overhead.
Datadog	<b>Datadog</b> is a full-stack cloud-based monitoring and alerting service for IT operations and development teams. It has containerized agents that can monitor container environments.
Dynatrace	<b>Dynatrace</b> has a new suite of monitoring tools, based on its Ruxit technology, for container monitoring and alerting. Its agent is injected into a container, where it can then auto-discover new services running on a host and fetch data from the Docker API. Dynatrace is also developing artificial intelligence technology to help with root cause analysis.
Elastic	<b>Beats</b> are single-purpose data shippers that install as lightweight agents and send data from machines to Logstash or Elasticsearch. Dockbeat is the shipper for Docker containers.
InfluxData	<b>Telegraf</b> is an open source agent that collects metrics with a Docker plugin; it's part of the stack that <b>InfluxData</b> is creating to collect and analyze metrics. The metrics are sent to InfluxDB or a variety of outputs.
New Relic	<b>New Relic</b> has updated its suite of products around their Digital Intelligence Platform for application and infrastructure performance. Its agent-based approach is particularly good for troubleshooting code-related application performance issues.
Sematext	<b>Sematext</b> has a Docker-native monitoring and log collection agent that collects and processes Docker metrics, events and logs for all cluster nodes and all auto-discovered containers. It has its own log management solution as well as performance monitoring service.
SignalFx	<b>SignalFx</b> has a collectd-docker plugin that captures metrics about CPU, memory, network, and disk using Docker's stats API. SignalFx has built-in Docker monitoring dashboards, which let you add dimensions (tags) metadata to aggregate, filter, and group metrics by property.

**TABLE 1:** There are a variety of other tools and solutions that can be used for identifying and collecting container data.

# Conclusion

All of the tools we've covered will be able to get you started with container visibility and monitoring. It's up to you to decide what combination of tools is needed to collect the information that will help you monitor and optimize your applications and systems. We covered the basics that everyone needs to know, which includes using the native functionality of Docker's Stats API and cAdvisor. That will go a long way towards introducing users to the basic necessities of monitoring containers. We also covered newer solutions such as Prometheus and Sysdig.

There are many ways to monitor containers, and solution vendors have opinionated ways of solving these problems. You might see this as a choice between open source and SaaS solutions, but it has more to do with finding the combination of solutions that fits your workload. No matter how you ultimately build your monitoring stack, you still need to know how to collect the metrics you're looking to monitor.

# THE IMPORTANCE OF HAVING VISIBILITY INTO CONTAINERS




In this discussion with Loris Degioanni of Sysdig, we talk about how the speed at which containerized environments move has proven to be difficult for traditional monitoring solutions. Users may have tens of thousands of containers, with some lasting only a few seconds, and a traditional monitoring solution whose sampling and reaction time is measured in minutes. It's important to have the contextual information about these containers in order to handle them in an environment where orchestrators are constantly spinning up and removing containers. Without context, you can be quickly overwhelmed by metrics that aren't useful to you. Degioanni talks about how Sysdig achieves visibility with its own ContainerVision technology. Instead of placing an agent inside the container, ContainerVision places a module in the operating system, allowing it visibility into all current and future containers. We also go on to talk about the container orchestration revolution, overcoming telemetry issues, and updating traditional monitoring solutions. [Listen on SoundCloud](#) or [Listen on YouTube](#)



**Loris Degioanni** is the CEO and founder of Sysdig, the Docker monitoring platform. He is the creator of the source troubleshooting tool Sysdig, and the open source container security tool Falco. Prior to founding Sysdig, Degioanni co-created Wireshark, the open source network analyzer, which today has more than 20 million users. Degioanni holds a PhD in computer engineering from Politecnico di Torino and lives in Davis, California.

# THE RIGHT TOOL FOR THE JOB: PICKING A MONITORING SOLUTION

by **BRIAN CHRISTNER**

onitoring is the hidden hero inside container infrastructure. Much like driving a car, we don't think about what's happening under the hood until something goes wrong. We want to pivot away from being reactive with our infrastructure and container problems, and move to a more proactive state of operation using monitoring insights to keep us up and running.

Gone are the days of just monitoring physical servers and their resources. We now need to know which containers are running and where they're running, monitor applications and containers, and provide alerting and metrics information to operations teams. The following sections will assist you in navigating the business decisions required to find the right monitoring tool for the job.

## Build vs. Buy

When considering new software, the decision to build or buy your solution should be at the top of the list. This argument is very similar to building a house from scratch versus buying an existing, finished house.

Building a house allows you to make changes and improvements as you perform the build, compared to buying a house where the decisions have mostly been made for you. Building is resource intensive, as it requires testing different components, trial and error experiences and learning how to handle new tools.

Buying a monitoring solution offers a finished product that's ready to run. This approach can save time and money, but there can also be drawbacks. You need to consider the anticipated growth of your organization. A rapidly changing environment with many new services being added can require significant additional time and effort to calibrate the monitoring solution. If many more people are expected to need access to the monitoring solution, then difficulties with identity management could be costly if the vendor charges for seat licenses.

Many monitoring solutions charge per agent. A solution's pricing model can be a deal-breaker depending on the size of your current or future deployment. For example, some systems charge per agent, while we have seen a few that actually charge per container monitored. Others are usage pricing based; for example, based on the data ingest rate of the agent.

Whether you are building or buying, vetting is required. While a mature product will have reduced risk, rapidly changing ones often have the latest features. An active community of users is important, whether you are going the vendor route or building something based on open source projects. The community is often responsible for keeping documentation and code updated.

Buying an off-the-shelf solution alleviates a lot of these questions, but a comparison between solutions should be performed to see exactly what the support provides, what service-level agreements (SLAs) are available and if professional services can assist with installation and customization.

## Support and SLAs

Support is a critical component to any monitoring system. Self-hosted solutions require maintenance, but you may not have the in-house operational capacity to build and manage a monitoring system on your own. But more to the point: who is responsible for supporting the monitoring solution? It is also important to ensure your company SLA and IT governance are adhered to when vetting a monitoring solution.

## Flexibility and Customization

Having the ability to customize a monitoring stack exactly to your requirements is an enormous advantage, but it comes at a price, and that price is the time required to make these customizations.

The build approach offers more flexibility and customization as you are working with open source software (OSS) that allows you to make your own changes. Buying offers some flexibility and customization dependant on the vendor's offering, but is relatively more restrictive compared to OSS alternatives.

## Cloud-Based vs Hosted On-Premises

Many monitoring solutions can run either in the cloud or be hosted on-premises. Each offers their own advantages and disadvantages, as well as several considerations based on your workload. Cloud generally means less control of your workload, whereas on-premises offers more fine-grained control. You should also consider if your monitoring needs have to adhere to regulatory or data storing requirements.

Generally, on-premises costs more because of the higher operational and support costs associated with running a datacenter, whereas cloud-based providers handle this for you when deciding for a cloud solution.

## Cloud

Moving to the cloud from on-premises is a [business case](#) in itself. Cloud offerings enable reduced costs by sharing resources with other cloud customers. The ability to quickly scale demand and the need for fewer specialized employees are all benefits related to cloud-based services. Central management of your infrastructure, including monitoring, is made easier with cloud vendors like [Amazon Web Services](#) (AWS), [Google Cloud Platform](#) (GCP), [IBM Bluemix](#) and [Microsoft Azure](#).

Several cloud vendors also offer one-click deployment of monitoring solutions, which allows you to get quickly up and running with a Software as a Service (SaaS) or open source solution. However, you must consider where the cloud is located compared to the your workload. Some users will want to have their monitoring stack located as close to their container workload as possible.

## Hosted On-Premises

Not everything can run in the cloud for a variety of reasons, including data privacy, licensing, internet connectivity and performance. On-premises solutions allow you to be in complete control of your environment. You're able to add more bandwidth and use custom hardware configurations. On-premises is sometimes combined with cloud services to offer hybrid solutions, which offer the best of both worlds. Additionally, not all SaaS monitoring solutions offer an on-premises solution, so ask about it when vetting different vendors.

There are other valid reasons for looking at hosting your own monitoring solution. The golden rule of going down the self-hosted route is that, if you are relying on alerting from your monitoring solution to notify you of critical incidents, then never host it in or on the same infrastructure as your application. That may seem obvious, but you will be surprised at the

amount of times it happens. If you take the cloud route, then that shouldn't be a problem.

The on-premises route offers a lot of flexibility and configuration possibilities, but will be more challenging to install and configure. The time it will take to have a monitoring solution up and running will take considerably more time compared to a SaaS solution that could be deployed instantly. These are all considerations that should be made when considering an on-premises solution.

## Compliance and Data Protection

Something that is often overlooked when considering a monitoring solution is the regulatory implications of data handling. Can data be stored outside of your data center, city, region or country? In some cases, depending on what industry you are working in, the government may also have regulations on how the data is handled.

If you install a cloud monitoring agent on a host machine that is shipping both metrics and information on what processes are running, then data ownership can become an issue. If your metrics include customer data, make sure your cloud provider can remove user data from the transactions it captures.

## Integration

It's important to fully understand what is currently running in your environment. Based on what you're running, you can start answering questions like what programming languages, applications and tools are currently or planned to run in your containers. This information will help you build a requirements list. This list will assist you in choosing what client libraries and integrations are important to your company now and in the future.

## Orchestrators

It is important to consider how orchestrators are monitored and to what extent. Consider the discovery of new services as they are brought online and what configurations are required in order to monitor these new services. Some monitoring tools will auto-discover new services or orchestrator nodes as they go online or offline, while other tools require configuration or integration with each service or node. The more auto-discovery capabilities available within the monitoring system, the less operation support required.

## Client Libraries and Integrations

Client libraries allow developers to write internal metrics within an application and expose them directly to a monitoring solution. Monitoring systems also integrate directly to the more common applications like databases and proxies. These integrations allow you to easily implement your monitoring solution and gather application-specific information. For example, [HAProxy](#) can be integrated to pull all the HTTP traffic statistics and display this information in monitoring dashboards. However, some integrations require additional agents, containers or sensors to collect metrics from the different integrations. Depending on the integration, this could range from easy to very complicated, and may require additional configuration and testing.

## Metric Consolidation

If you are currently using a cloud provider or Platform as a Service (PaaS), chances are they're already exposing monitoring metrics. Find out what is available within your current Containers as a Service (CaaS), PaaS or Infrastructure as a Service (IaaS) stack that you can leverage. Can you integrate metrics from the cloud service stack into your monitoring solution to create a single view of infrastructure and containers? You may want to centralize metrics into a single solution rather than using many

different tools without integrations. Be sure to find out which of your current services can be repurposed into your new monitoring solution.

## Conclusion

We have explained the use cases for build versus buy, cloud versus on-premises, integrations and native functionality for monitoring solutions. Based on this information, you can start building a matrix of information and make a comparison between the different monitoring solutions currently available on the market, both open source and SaaS solutions. Do your homework and gather information internally first to understand your requirements. This requirements list will be your key to choosing the right monitoring tool for the job.

# UNDERSTANDING THE APPLICATION PATTERN FOR EFFECTIVE MONITORING



In this conversation with Ken Owens, chief technology officer, Cloud Native Platforms at Cisco Systems, we talk about what makes monitoring containers different, especially when it comes to

behaviors like networking. Containers are a small unit of capability that result from when you decompose an application into a set of services. Understanding the availability, liability and the translation to these different small microservices is much more complicated and difficult to do than in the past. Owens also talks about the kind of complexity that the microservices pattern introduces. In an environment where containers are being automated and orchestrated at scale, it becomes difficult to identify and isolate problems. However, if you understand the application pattern you're trying to achieve, you know how to monitor that pattern.


[Listen on SoundCloud](#) or [Listen on YouTube](#)



**Ken Owens** is chief technology officer, Cloud Native Platforms at Cisco Systems. Owens is responsible for creating and communicating both technical and scientific vision and strategy for CIS business. He brings a compelling view of the technology trends in enterprise IT (e.g., infrastructure, computing, SaaS, virtualization and cloud) and evangelizes the technology roadmap for the business. Before joining Cisco in 2014, Owens spent over seven years at Savvis as the chief scientist, chief technical officer and vice president of Security and Virtualization Technologies.

# MANAGING INTENTIONAL CHAOS: SUCCEEDING WITH CONTAINERIZED APPS

by **VIVEK JUNEJA**

hile building applications that use containers has become an established practice, some organizations are facing challenges with managing containers. It's no longer rare for engineering teams to have working knowledge of using containers in their routine development workflow. However, managing containers in production environments requires an even greater understanding of container technology, along with unlearning habits and practices that no longer work. These practices span across the different life cycle aspects of container management: logging, monitoring, notifications, remediation and metrics. Missing out on any one of them could shake a team or organization's confidence in containers.

Containerized applications have to be designed with consideration for the various nonfunctional and container-native patterns that have emerged over time. For example, the [twelve-factor application style](#) has been widely advocated and adopted by microservices enthusiasts and developers. The twelve-factor principles requires the application to be operation-ready, including design for fast startups and graceful

shutdowns, but it misses out on handling the [software entropy](#) that such container-based systems acquire over time.

## Containers Add Complexity to the Overall System

Adopting containers, especially in a microservices-based architecture, leads to increased entropy in the system. When we say entropy, what we mean is the system becomes immensely more complex, and that complexity needs to be managed. The complexity is due to the increase in moving parts in a container-based environment. Moreover, the disposable and immutable nature of containers can encourage some teams to create more rapid delivery pipelines. This propensity towards change further fuels the overall system complexity. If not managed, it can create opportunities for burnout, downtime, and disappointment in adopted technologies.

Container adoption has led to the development of new practices from web-scale organizations. The major lesson from public cloud providers — advocating the need to design for failure and redundancy at all levels — is finally going mainstream. Another important lesson is to have applications that are designed to handle the complexity from containers, a lesson that might remind many of the [antifragile movement](#).

## Understanding New Design Goals

In order to address these challenges with containerized applications, there are ways to create containerized applications in a way that makes them operation-ready. Making changes to the development process will enable better monitoring and management of containers in production environments. These lessons translate into some high-level design goals for containerized application development.

## 1. Applications running as containers should acknowledge that services could be interrupted at any time.

Interruption could be triggered by a variety of situations. In the case of Docker, a stop command sends the SIGTERM signal to the application, indicating the request to shut down. This may require applications to perform cleaning-up activity before the container dies. Taking a hint from the twelve-factor principles, the application must gracefully shutdown. For example, if there is a state that is being manipulated by the container, that would require being checkpointed. The [checkpoint](#) would be persisted onto external storage, where other container workers could access it.

## 2. An application should expose an interface to allow for health checks that can be used by the container platform.

The accuracy of the health check implementation is critical to the functioning of this setup. Usually, an application emits health check information that indicates it can correctly connect to third party or external interfaces. A failed health check is a good indicator for the platform to remove the instance from the configuration. Platforms also use this to decide remediation activities that need to be performed due to an error in the container. This remediation could involve launching the same container image on a different host.

## 3. There should be a mechanism to reliably identify a faulty instance and its root cause.

Diagnosing application issues in a containerized environment is an open invitation to chaos. The problem occurs when there is a need to identify the exact container instance that caused the issue. Using a common identification artifact that spans across the notification, monitoring and logging infrastructure is one way to solve this problem. A common approach is to use a business-specific identity, configured

to be emitted as a notification, that can be captured by the logging system. The log should then provide mapping between the identifier and the container name that caused the issue.

#### 4. **The log format must provide full context to an operator.**

Logging remains an important element in the arsenal for debugging issues in the production environment. With container environments, the log format must include some context like container ID, container image and container host. The container runtime injects a HOST environment variable that can be used by the application in its logs. Container platforms use metadata in their log stream, which helps identification and remediation activities.

## Making Agile Decisions with Containers

A fast and agile environment provides the ability to make quick decisions and iterate as needed. One of the cornerstones of the DevOps movement is the adoption of ideas around the observe, orient, decide and act ([OODA loop](#)), a widely-used principle to take action in a constantly changing environment. Mapping the OODA practice to a containerized production environment would lead to the following inferences:

- **Observe:** This pertains to alerts and notifications that filter out the useful signals from the noise. This is possible through tools that receive events from the monitoring system when something goes wrong. Having a good signal-to-noise ratio at this stage is critical to the overall success of the process.
- **Orient:** Once access to the information is sorted, it is used to identify the symptoms causing an issue. Getting information from the logging and monitoring system is the basis of orientation. You must be able to identify the exact source of information with minimal noise.

- **Decide:** Based on the symptoms identified during the orientation phase, you must decide what action to take to resolve the situation. An example action would be changing the group configuration or relocating to a new set of hosts. If the issue identified is related to the application logic, then rolling back to the previous configuration could be a possible fix.
- **Act:** The container platform and tools must allow for fast action once they're decided upon. Having access and permission to the container management tools is useful.

Container implementations in the enterprise must allow the OODA loop to be implemented and have fast transitions. The merit of any container management and monitoring system is measured by the accuracy of the information it provides and the speed with which actions can be taken.

## Tools for Taming Complexity

The art of managing the chaos in a container-based environment has led to the creation of new tools and services that embrace the dynamic and autonomous nature of container management. Tools like [StackStorm](#) and [Netflix Winston](#) have inspired implementations to trigger automated workflows in the case of events, especially events that involve an issue with the environment. Tying this to your container platform can allow operation runbooks to execute in case of a fault with a container. This reduces manual intervention and engineering team burnout, which increases productivity.

One of the concepts we discussed earlier was to monitor groups of containers instead of focusing on individual instances. The use of container labels and environment variables can be used to implement this practice. A tool like [cAdvisor](#) can capture the labels provided to any

container on a host. If environment variables are used, cAdvisor also allows them to be captured using the --docker-env-metadata-whitelist runtime parameter.

Tracing calls between self-contained services in an architecture is difficult with traditional practices. Improving practices around tracing is an important part of continued success with microservices. Tracing platforms like [OpenTracing](#) will become commonplace in all container based environments going forward.

The Cloud Native Computing Foundation has adopted OpenTracing as a hosted project. There are also tools like [Zipkin](#), an [open source tracer for microservices](#), first developed by Twitter to track web requests as they bounced around different servers. There's also [Sysdig tracers](#), which allows for open source tracing of everything from microservices down to system resource access time.

Taking actions iteratively in an OODA loop is an important part of container implementation. Platforms like [Vamp](#) allow workflows to be implemented for making canary release decisions based on containerized application metrics. Tools like this could act as a method of implementing the OODA loop and applying it to release and testing practices.

If you are running containers with process isolation, then finding a flagged process running among a set of containers across hosts is a challenging feat. Identifying the host that runs the flagged container is one part of the problem. Usually, this is solved through a host monitoring tool like [Prometheus](#). Once you identify the host, you can perform process ID mapping between the host and the container. This requires identifying the root process ID and correlating it with the running container. Tools like [Sysdig](#) solve this problem and much more with little or no overhead on the container performance.

[Cloud Foundry](#) has a unique approach to solving container management and monitoring difficulties. It provides an abstraction of the containerized infrastructure in the form of well-designed APIs, a solid development environment, and by providing logs and metrics for each application. These features make it easy for developers to adopt agile practices and leverage visibility into their containerized production applications.

## Conclusion

Organizations working in a hybrid setup, involving both containerized and traditional workloads, will have a hard time embracing this shift. The challenge is maintaining systems that revolve around different schools of thoughts. Legacy systems are usually long running and non-disposable; they demand per instance, granular approaches to monitoring and management.


Some organizations will want to experiment with a team that is independent and has few traditional system monitoring needs. However, the OODA loop is still a valuable approach to containerized applications, and establishes common ground rules for both traditional and container-based environments.

Developers need to be more aware of new practices in monitoring, logging and service management in order to create containerized applications that will thrive. The changes needed to successfully adopt containers will result in cloud-native systems that are able to accomplish the goals of agile IT.

*This is a case study provided by our ebook sponsor, Sysdig.*

# MONITORING DOCKER AND KUBERNETES AT WAYBLAZER

by **APURVA DAVE**

oving Docker into production is still as much art as it is science. We caught up with Kevin Kuhl and Kevin Cashman at WayBlazer to talk about monitoring Docker and Kubernetes, selecting the right metrics to focus on, deciding what to run with Kubernetes and diagnosing the “domino effect” in a Kubernetes failover.

## Introducing WayBlazer

[WayBlazer](#) is the world’s first cognitive recommendation engine focused on delivering artificial intelligence to the travel industry. Bringing together IBM Watson and proprietary cognitive computing technology, WayBlazer’s AI analyzes cues and triggers from the traveler’s search to yield personalized hotel recommendations, insights, images and local area knowledge in the context of each individual and their trip profile. The combination of personalized content with hotel recommendations boosts online engagement and improves traveler conversion rates.

Based in Austin, Texas, the company is led by seasoned travel and tech entrepreneurs, including the founder of Travelocity and founding chairman

of Kayak.com, the former general manager of IBM Watson, and the former president of Sabre Hospitality Solutions.

## What's your environment like?

We're an AWS shop, so as you can imagine we heavily use EC2. We also take advantage of a number of AWS services; they're good, cost effective, and managed really well, so that makes our lives easier. Not too long ago, we started using Kubernetes, and now we run around a dozen Kubernetes minions, sometimes more when we scale, on top of our EC2 nodes. We run about 15-20 containers per node. In addition, we have a number of Kubernetes namespaces – prod, test, and project-specific namespaces.

Within that environment, we're running a number of different services. Our frontend is in JavaScript and our API service is running on Node.js. We have Java services doing the heavy lifting in the backend, as well as some auxiliary applications running in Python. We use Sysdig and AWS Cloudwatch to provide monitoring.

## Do you run everything in Kubernetes?

No, not everything. Our approach is to start by saying, “should it run in Kubernetes?” We have a number of areas where the answer is no. If we don't run something in Kubernetes either it's not a functional fit or we didn't like the performance profile.

The most obvious things we run inside it are our stateless services. They're the perfect fit. Stateful services, like our graph database, aren't a good fit today. Maybe with PetSets we'll eventually be able to do that, but we don't right now. Finally, we have a lot of jobs we run once or we run on a schedule. We prefer to run these outside of Kubernetes. Again, it's something that Kubernetes is building towards but doesn't meet our needs today.

## How do you keep track of everything you're monitoring?

Well, way back, we relied solely on AWS Cloudwatch to give us host-level metrics. That worked OK before we had Docker and Kubernetes, as we were operating a more static architecture. Also, when things were less dynamic you could derive a lot more information about the application's performance from your infrastructure metrics. Moving to Kubernetes changed that for us. We needed to have container and service level monitoring.

Kubernetes came packaged with Heapster and Grafana, so that was the first thing we tried. It was a very heavy investment in our time and effort to get it running and then to manage it. Given we're an ops team of two, it wasn't the most strategic way to spend our time. We wanted a monitoring capability that was full-featured, robust, and we could hit the ground running with. So for us, that ruled out the roll-your-own, glue-some-components-together approach. The choices we came down to were commercial products that would fit our needs. We actually still use CloudWatch for basic AWS metrics, and now also use Sysdig for all the higher-level stuff. There's a bit of duplication there in terms of monitoring low-level infrastructure, but that's OK by us.

## Does having Kubernetes impact how you monitor Docker?

Yes, but Kubernetes has impacted a lot of things, not just monitoring. If we go back to why we chose Kubernetes, one of the things we liked was that it precisely described what it was doing and had its own language to describe it: namespaces, deployments and so on. That "language" is opinionated in a way that forced everyone here to get on the same page when it comes to our container technology. That made it easier to communicate with engineers, ops... basically anyone involved with the

creation and operation of our software. Part of DevOps is putting your infrastructure metrics in front of engineering — it was essential to have them consistent.

To get more specific on monitoring, by far the biggest, most obvious thing we realized moving to Kubernetes was that we needed monitoring capabilities that could give us insight into a Kubernetes-centric and a service-centric view of our code. Let's face it, a lot of things can monitor Docker containers at the infrastructure level by talking to the stats API, but they don't know anything about what that container is.

Knowing that these 10 containers make up the API server, or this one container is part of this pod, all that knowledge is hugely valuable and has reshaped monitoring for us.

## **What are the most important monitoring metrics for you? What is it that you track day in and day out?**

The funny part is, we're talking about Docker monitoring, but we don't really monitor Docker much at all. We focus on monitoring our applications, and to some extent we monitor Kubernetes to give us environmental information about our applications.

We will often use Docker related information for troubleshooting, it's another layer in our infrastructure stack, and may often present information that helps us solve tricky problems. Here's some of the key stuff we monitor at the application level:

- **Request times:** We keep a close eye on the request and response times for our frontend and API. These are the first indicators that will fire if something goes wrong or changes. For example a new build could throw our application off, or a network misconfiguration could have a negative effect on the user.

- **Backend query response time:** These could be directly related to a database or to our Java backend. Again, this is a good place to look for potential problems brewing, as well as for opportunities to improve our system. If a commonly executed query is running slowly, we can bring that back to the team for investigation and remediation.
- **Heap usage and garbage collection:** While we don't use them regularly, we've used this data to debug our data warehouse based on JVM-level data. Again, we're really looking at these resource metrics from an app level.

Next is Kubernetes. Here's what we track:


- **Pod restarts and container counts:** They tell us if something is changing.
- **Failed launches:** These are important events that we can use to correlate with other metrics, whether app or infrastructure.
- **Service request latency:** This one is challenging. It involves aggregating the latency information from all our containers from a particular service and presenting that in one, unified view. This is a prime example of having a “Kubernetes viewpoint” in your monitoring.
- **Per-container resource utilization:** We've recently added requests and limits in Kubernetes, which means we can track resource utilization more effectively. Instead of tracking CPU and memory as a function of the host, we track these metrics based on the allocated resources to that particular container. That should give us an earlier warning to resource allocation issues as opposed to just monitoring at the host utilization level.

Yes, we also have typical alerts on high CPU, memory, and disk utilization on critical nodes: you have to have them.

## Can you tell your story about monitoring the Kubernetes “domino effect”?


This was something we experienced early on with running Kubernetes in EC2 as we tested our new environment out. Our first warning indicator was an alert around high API latency. At this point latency started to subside – we were ready to act if needed but things seemed to have settled down.

But pretty soon we knew there was a serious problem. We started seeing alerts in Amazon for unhealthy instances autoscaling group for Kubernetes manages. We immediately went to Sysdig and noticed the blue line going down. This wasn't good – but not yet the end of the world. Kubernetes was acting by failing those containers over to another node.



**FIG 1:** Kubernetes migrating failed containers to another host.

Then saw both disk usage and memory usage go very high at the same time. This was the final death knell. This entire system fell over.



**FIG 2:** A domino effect of failovers.

Essentially we had a domino effect – Kubernetes failover worked on the first node, which had purged all of the containers when it ran out of memory. At this time requests and limits for memory usage were not present on some services.

The trouble was that the other nodes in the cluster already had enough disk usage that they couldn't reserve enough space to fully migrate that number of containers to other nodes.

We were quickly able to widen the autoscaling group and heal a couple nodes to get us back on track; but it was painful and not ideal. During our retrospective on this, it was interesting to be able to review the whole thing through Sysdig. We very quickly were able to have a very robust report with the team within a day showing exactly what happened. We generated a full list of action items with specific numbers for thresholds, limits, scaling policies and requirements. It generated several sprints of production hardening work and made our systems much more robust.

## Summary of Lessons Learned

Our experience shows that adopting Kubernetes and Docker will have a significant impact on your monitoring approach. Much like Kubernetes can move you to a more service-centric view of your infrastructure, your monitoring strategy will move the same way. You're going to need both an infrastructure-based view and a service-centric view in your monitoring system to run containers in production.

Make sure your monitoring system reinforces the common language of Kubernetes. It must understand the basic constructs of Kubernetes out-of-the box. You shouldn't have to teach it anything. Once you deploy it and instrument your hosts, you should get a service-level view of your metrics. The next big thing is that your instrumentation should see inside

containers to get a more application or code-centric view of our applications. Although we keep tabs on infrastructure metrics, we focus on monitoring our application, even as Kubernetes is scaling Docker containers up and down.

All of these aspects reduce cost and make us more efficient. It means we don't have to put engineering effort into adapting raw metrics to reflect that Kubernetes metadata, nor do we have to think too hard about instrumenting our systems.

# WHAT ARE YOUR DOCKER CONTAINERS DOING?



**“Sysdig gives my team unprecedented visibility into our applications. We're using it extensively to help us solve complex issues in our most innovative products”** – Adam Hertz, VP of Engineering at Comcast

*We're about to make your life easier.*

Enterprises are moving to Docker containers for faster software development that promises developer agility, software portability, and scalable microservices. That's all good news.

Here's the bad news: Docker environments get harder to operate, because microservices and containers break legacy monitoring tools. Containers are great for developers, but they will wreak havoc on your DevOps team.

*Worried? Relax, we got your back.*

Sysdig is the first and only solution that can natively monitor Docker environments. Sysdig ContainerVision™ provides request-level visibility inside containers without invasive instrumentation. This approach overcomes the limitations of your old monitoring tools, and at the same time makes monitoring Docker simpler and more robust. All that means you can sleep at night when Kubernetes is auto-scaling your apps. Go on, press the snooze button.

Visit [sysdig.com](http://sysdig.com) and find out!

# CONTAINER MONITORING DIRECTORY

Although this directory has over a hundred entries, it is not meant to be comprehensive. Instead, it lists many of the projects and vendor offerings that are used to monitor containers. Listings are divided into four sections to make it easier for the reader to quickly review. Since most of the products actually provide a range of functionality listed, please use the categorization as a starting point when reviewing solutions you may want to use or consider.

# CONTAINER-RELATED MONITORING

*This section includes tools that are purpose-built to monitor containers. In addition, it includes commonly used monitoring tools that have extensions or plugins to collect and monitor container data.*

Product/Project (Company or Supporting Org.)	Functionality Includes:
<b><a href="#">AppDynamics (AppDynamics)</a></b> Application and business performance software that collects data from agents installed on the host. It provides an extension to collect data from the Docker API.	» Performance Monitoring
<b><a href="#">AppFormix (AppFormix)</a></b> Cloud infrastructure monitoring and analysis software that runs in any public, private, multi-tenant, or hybrid environment. It includes ContainerFlow, which utilizes Intel Resource Director technology to enforce isolation between workloads and deliver better performance for container workloads. The company is focused on analytics for operators of OpenStack and Kubernetes.	
<b><a href="#">Bluemix Container Monitoring and Logging (IBM)</a></b> Monitoring and logging is a built-in service for Bluemix Containers, integrated to the cloud offering. Users receive monitoring information and logging capabilities for any container they deploy on Bluemix without having to configure or install anything. This service is based on Agentless System Crawler for data collection; it can utilize Logstash, Graphite, Elasticsearch, Kibana and Grafana for data ingestion, analysis and visualization.	
Open Source <b><a href="#">Cabin (SkippBox)</a></b> An iOS application for managing Kubernetes applications. An Android application is being developed.	» Management/Orchestration
Open Source <b><a href="#">cAdvisor (Google)</a></b> cAdvisor (Container Advisor) is a Google-supported project that analyzes resource usage and performance characteristics of running containers.	
<b><a href="#">Cobe.io (Cobe.io)</a></b> Provides a live topology of heterogeneous infrastructure, on top of which model performance metrics and alerts are overlaid.	
<b><a href="#">ContainerX (Cisco)</a></b> ContainerX is in the process of being integrated into Cisco's Cloud group. It provides a single pane of glass for containers running on bare metal or VM, Linux or Windows, private or public cloud. It provides APIs for enterprises wanting to integrate the platform's capabilities into their custom management portals.	» Management/Orchestration

	Product/Project (Company or Supporting Org.)	Functionality Includes:
	<b>CoScale (CoScale)</b> Combines server and container resource metrics with application-specific performance metrics. Its lightweight agents have little overhead.	» Anomaly Detection
Open Source	<b>DC/OS Metrics (Mesosphere)</b> An experimental project to make metrics more accessible. It is expected to be incorporated in a future release of DC/OS 1.10.	
Open Source	<b>Dockbeats (Ingensi)</b> Dockbeat is a lightweight agent that is installed on servers, which periodically reads Docker container statistics and indexes them in Elasticsearch. It is based on the Elasticsearch Beats model.	
	<b>Docker Datacenter (Docker)</b> Provides on-premises container management and deployment services to enterprises with a production-ready platform supported by Docker and hosted locally behind the firewall. The included Universal Control Plane enables teams to manage and deploy Dockerized applications, and can be run on any private infrastructure or public cloud.	» Management/Orchestration
Open Source	<b>Docker Engine (Docker)</b> A lightweight runtime and tooling that builds and runs your Docker containers. The Docker Stats API in Docker Engine is the source of most information collected by tools that monitor containers. CS Docker Engine is the commercially supported version. Docker Content Trust is a feature that makes it possible to verify the publisher of Docker images.	
	<b>Dynatrace (Dynatrace)</b> Dynatrace's new suite of monitoring tools is based on its Ruxit technology. Its agent is injected into a container, which then autodiscovers new services running on a host and can fetch data from the Docker API. Dynatrace is also developing artificial intelligence technology to help with root cause analysis.	» Performance Monitoring
Open Source	<b>Falco (Sysdig)</b> A behavioral activity monitor designed to detect anomalous activity in your applications. Powered by Sysdig's system call capture infrastructure, Falco lets you continuously monitor and detect container, application, host and network activity.	» Anomaly Detection
Open Source	<b>Heapster (Cloud Native Computing Foundation)</b> Enables analysis of compute resource usage and monitoring of container clusters. Heapster currently supports Kubernetes and CoreOS natively.	
	<b>HPE SiteScope (Hewlett-Packard Enterprise)</b> Agentless monitoring software for monitoring the availability and performance of distributed applications and IT infrastructure. The dynamic Docker monitor of SiteScope monitors the health and performance from the cluster manager level through the Docker daemon layer to the container level, including the applications running inside the containers.	
Open Source	<b>Kubernetes Dashboard (Cloud Native Computing Foundation)</b> A general purpose, web-based UI for Kubernetes clusters. It allows users to manage applications running in the cluster and troubleshoot them, as well as manage the cluster itself.	

	Product/Project (Company or Supporting Org.)	Functionality Includes:
	<b><a href="#">Operations Management Suite (Microsoft)</a></b>	
	Created as a management and security tool for hybrid clouds. It also provides a simplified view of containers' usage and diagnoses issues.	
Open Source	<b><a href="#">Prometheus (Cloud Native Computing Foundation)</a></b>	» Time Series Database An open source, service monitoring system and time series database.
	<b><a href="#">Scout (Scout)</a></b>	» Performance Monitoring A Docker monitoring partner, Scout provides application and server monitoring, especially for Rails applications. It has plugins to monitor both Docker and LXC containers.
Open Source	<b><a href="#">Sematext Docker Agent (Sematext)</a></b>	It collects and processes Docker metrics, events and logs for all cluster nodes and all auto-discovered containers. It works with Docker Swarm, Docker Datacenter, Docker Cloud, as well as Amazon EC2, Google Container Engine, Kubernetes, Mesos, RancherOS and CoreOS.
	<b><a href="#">SignalFx (SignalFx)</a></b>	» Time Series Database Its collectd-docker plugin capture metrics about CPU, memory, network, and disk using Docker's stats API. SignalFx has built-in Docker monitoring dashboards, which let you add dimensions (tags) metadata to aggregate, filter, and group metrics by property.
	<b><a href="#">SPM Performance Monitoring (Sematext)</a></b>	» Anomaly Detection » Performance Monitoring On-premises performance monitoring, alerting and anomaly detection. An agent is used to monitor Docker containers and host metrics about CPU, memory, network, disk and events.
Open Source	<b><a href="#">Sysdig (Sysdig)</a></b>	Linux system exploration and troubleshooting tool with first class support for containers. This is the open source project that Sysdig Cloud is based on.
	<b><a href="#">Sysdig Cloud (Sysdig)</a></b>	Based on open source Sysdig technology, Sysdig Cloud monitors, troubleshoots and alerts on containerized environments. Sysdig Cloud can be used as a cloud service or deployed as hosted software in your private cloud.
Open Source	<b><a href="#">Weave Context (Weaveworks)</a></b>	API-compatible Prometheus implementation that natively supports multitenancy and horizontal scale-out clustering.
	<b><a href="#">Weave Scope (Weaveworks)</a></b>	Weave Scope offers a real-time monitoring solution for containers.
Open Source	<b><a href="#">Zenoss Control Center (Zenoss)</a></b>	Zenoss Control Center is an application management and orchestration system. It works with the Zenoss platform and other Docker applications. Serviced is a popular repository in this project that provides a PaaS runtime.

# COMPONENTS/CLASSES OF MONITORING SYSTEMS

*A full stack monitoring system has many different components. This section includes tools that collect and ingest, process and store data and metrics. It also includes the processing of events and the creation of alerts. Other important parts of a monitoring system are the visualization and tools to help with tracing, error tracking and debugging.*

	Product/Project (Company or Supporting Org.)	Functionality Includes:
Open Source	<a href="#">Agentless System Crawler (IBM)</a> A unified cloud monitoring and analytics framework that enables deep visibility into all types of cloud platforms and runtimes. It is used as the data collection framework for Bluemix Containers and Vulnerability Advisor.	» Collect/Ingest
	<a href="#">Beats (Elastic)</a> A platform for single-purpose data shippers. They install as lightweight agents and send data from machines to Logstash or Elasticsearch.	» Collect/Ingest
	<a href="#">Bosun (Stack Exchange)</a> A time series alerting framework developed by Stack Exchange. Scollector, a metric collection agent, is used with it.	» Alerting and Event Processing
	<a href="#">Circonus (Circonus)</a> Collects any data from any system, and then provides customizable visualization, analysis and alerting tools.	» Alerting and Event Processing » Time Series Database
Open Source	<a href="#">collectd (N/A)</a> A small daemon which periodically collects system information and provides mechanisms to store and monitor the values in a variety of ways.	» Collect/Ingest
Open Source	<a href="#">Fluentd (Cloud Native Computing Foundation)</a> Fluentd is an open source data collector for unified logging layers.	» Log Management
	<a href="#">Google Stackdriver (Google)</a> Monitoring, logging and diagnostics for applications on Google Cloud Platform and AWS.	» Log Management » Performance Monitoring
Open Source	<a href="#">Grafana (Raintank)</a> Grafana is a metrics dashboard and graph editor for Graphite, InfluxDB and OpenTSDB.	» Visualization

	<b>Product/Project (Company or Supporting Org.)</b>	<b>Functionality Includes:</b>
Open Source	<a href="#"><b>Graphite Carbon (N/A)</b></a> Component of the Graphite project that is responsible for receiving metrics over the network and writing them down to disk using a storage backend.	» Collect/Ingest
Open Source	<a href="#"><b>Graphite-Web (N/A)</b></a> Graphite-Web is a Django-based web application that renders real-time graphs and dashboards.	» Visualization
Open Source	<a href="#"><b>Graylog (Graylog)</b></a> Graylog is a fully integrated platform for collecting, indexing, and analyzing both structured and unstructured data from almost any source.	» Log Management
Open Source	<a href="#"><b>Hawkular (Red Hat)</b></a> Hawkular is a modular systems monitoring suite that consists of several subprojects for storing metrics, alerting on incoming events and more.	» Performance Monitoring
	<a href="#"><b>Honeycomb.io (Honeycomb.io)</b></a> Honeycomb is an event-driven observability tool for debugging systems, application code and databases. Honeycomb uses structured data and real-time query aggregation to support ultra-rich datasets without indexes or schemas, and a fast, interactive interface.	» Tracing/Error Tracking/Debugging
	<a href="#"><b>Incident.MOOG (Moogsoft)</b></a> Event management driven by real-time machine learning to detect anomalies across your production stack of applications, infrastructure and monitoring tools.	» Alerting and Event Processing
Open Source	<a href="#"><b>InfluxDB (InfluxData)</b></a> An open source database written in Go specifically to handle time series data with high availability and high-performance requirements. It is one of the more commonly used databases in the stacks used to monitor containers.	» Time Series Database
Open Source	<a href="#"><b>Kapacitor (InfluxData)</b></a> Open source framework for processing, monitoring, and alerting on time series data.	» Alerting and Event Processing
Open Source	<a href="#"><b>Kibana (Elastic)</b></a> Kibana is a flexible analytics and visualization platform.	» Visualization
	<a href="#"><b>Librato (SolarWinds)</b></a> Librato Agent is used to collect Docker metrics. Librato itself manages time series data, alerting and visualization.	» Alerting and Event Processing » Time Series Database
	<a href="#"><b>Lightstep (Lightstep)</b></a> A distributed tracing system built on top of the OpenTracing standard. LightStep collects, analyzes and presents high-fidelity end-to-end traces for the requests having the largest impact on your business, and the user experience of your system.	» Tracing/Error Tracking/Debugging
	<a href="#"><b>Logentries (Rapid7)</b></a> Logentries provides analytics tools to monitor Docker environments.	» Log Management

Product/Project (Company or Supporting Org.)		Functionality Includes:
	<b><a href="#">Loggly (Loggly)</a></b>	» Log Management Loggly collects log data using existing open standards like syslog and HTTP.
	<b><a href="#">Logsense (Semantext)</a></b>	» Log Management Logsense replaces Logstash in the ELK stack and is provided as a hosted solution.
Open Source	<b><a href="#">Logspout (Glider Labs)</a></b>	» Log Management A log router for Docker containers that runs inside Docker. It attaches to all containers on a host, then routes their logs wherever you want. It also has an extensible module system. It's a mostly stateless log appliance. It's not meant for managing log files or looking at history. It is just a means to get your logs out to live somewhere else, where they belong.
Open Source	<b><a href="#">Logstash (Elastic)</a></b>	» Log Management Logstash is a tool for managing events and logs.
	<b><a href="#">logz.io (logz.io)</a></b>	» Analytics » Log Management Log analytics data with hosted ELK as a service.
	<b><a href="#">Meros (Meros)</a></b>	» Log Management Meros offers a GUI that provides monitoring, management, log aggregation and alerts for Docker.
	<b><a href="#">NetViz (BanyanOps)</a></b>	» Visualization A platform to discover, visualize and analyze microservices and their interactions. Using lightweight, on-host, deep packet inspection technology, NetViz autodiscovers services and application resources and provides real-time network analysis and insights. Operations and security teams use NetViz to improve the performance, availability and security of their infrastructure.
Open Source	<b><a href="#">OpenTracing API (N/A)</a></b>	» Tracing/Error Tracking/Debugging Consistent, expressive, vendor-neutral APIs for distributed tracing and context propagation.
Open Source	<b><a href="#">Riemann (N/A)</a></b>	» Alerting and Event Processing A network event stream-processing system in Clojure that aggregates events from servers and applications.
Open Source	<b><a href="#">RRDtool (N/A)</a></b>	» Log Management » Time Series Database A high performance data logging and graphing system for time series data. RRDtool can be easily integrated in shell scripts, Perl, Python, Ruby, Lua or Tcl applications.
Open Source	<b><a href="#">Sentry (Functional Software, Inc.)</a></b>	» Log Management » Tracing/Error Tracking/Debugging Cross-platform error logging and aggregation.
Open Source	<b><a href="#">Snap (Intel)</a></b>	» Collect/Ingest An open telemetry framework designed to simplify the collection, processing and publishing of system data through a single API.

Product/Project (Company or Supporting Org.)	Functionality Includes:
<a href="#">Splunk Enterprise (Splunk)</a>	<ul style="list-style-type: none"> <li>» Log Management</li> </ul> <p>Simplifies the collection and analysis of data-generated technology infrastructure, security systems and business applications. Splunk offers a "solution" that focuses on container monitoring and application delivery. A cloud offering is also available.</p>
<a href="#">Stackify (Stackify)</a>	<ul style="list-style-type: none"> <li>» Log Management</li> <li>» Performance Monitoring</li> <li>» Tracing/Error Tracking/Debugging</li> </ul> <p>Provides software developer tools that fully integrate error and log management with application performance monitoring and management for server and developer workstations.</p>
<a href="#">Sumo Logic App for Docker (Sumo Logic)</a>	<ul style="list-style-type: none"> <li>» Log Management</li> </ul> <p>The Sumo Logic app for Docker uses a container that includes a collector and a script source to gather statistics and events from the Docker Remote API on each host. The app wraps events into JSON messages, then enumerates over all running containers and listens to the event stream. This essentially creates a log for container events. In addition to configuration information obtained using Docker's Inspect API, the app collects host and daemon logs. All this data is integrated into Sumo Logic's monitoring dashboard, which specializes in handling machine data analytics.</p>
Open Source <a href="#">Telegraf (InfluxData)</a>	<ul style="list-style-type: none"> <li>» Collect/Ingest</li> </ul> <p>An agent written in Go for collecting metrics from the system it's running on, or from other services, and writing them into InfluxDB or other outputs.</p>
Open Source <a href="#">Trace (RisingStack)</a>	<ul style="list-style-type: none"> <li>» Tracing/Error Tracking/Debugging</li> </ul> <p>A visualized stack-trace platform designed for microservices monitoring.</p>
<a href="#">Wavefront (Wavefront)</a>	<ul style="list-style-type: none"> <li>» Analytics</li> <li>» Collect/Ingest</li> </ul> <p>Uses cAdvisor to collect container metrics, which are analyzed along with metrics from other systems and applications.</p>
Open Source <a href="#">Whisper (N/A)</a>	<ul style="list-style-type: none"> <li>» Time Series Database</li> </ul> <p>A file-based time series database format for the Graphite project.</p>

# MANAGEMENT/ ORCHESTRATION

*Most container management platforms include dashboards that let operators monitor at least some metrics, often to enable the intelligent orchestration of services. This section includes only a sampling of available options.*

Product/Project (Company or Supporting Org.)	Functionality Includes:
<b>Apcera Platform (Apcera)</b> Apcera manages access to compute resources across a cluster of servers. By focusing on managing policies across multiple environments, it aims to secure workloads and containers in enterprise production environments. It provides cluster monitoring with Zabbix, which is preconfigured for monitoring Apcera Platform components.	» Management/Orchestration
<b>Azure Container Service (Microsoft)</b> Azure Container Service simplifies the creation and configuration of a cluster. The default configuration includes Docker and Docker Swarm for code portability, and Marathon, Chronos and Apache Mesos to ensure scalability.	» Management/Orchestration
<b>Cloudsoft Application Management Platform (Cloudsoft)</b> Cloudsoft's application management platform, based on the open source Apache Brooklyn project, orchestrates services, platforms and infrastructure, including deployment to containers.	» Management/Orchestration
<b>ContainerShip (ContainerShip)</b> ContainerShip is a self-hosted container management platform, capable of running on any cloud, and used to manage containers from development to production.	» Management/Orchestration
Open Source <b>Contiv Cluster (Cisco)</b> Contiv Cluster is an integrated collection of software that simplifies deployment, management and maintenance of clustering systems. It is part of Project Contiv, an open source project that defines infrastructure operational policies for container-based application deployment.	» Management/Orchestration
<b>EC2 Container Service (Amazon Web Services)</b> Amazon EC2 Container Service helps companies manage clusters of containers on AWS infrastructure.	» Management/Orchestration
<b>Elastic Cloud (Elastic)</b> An Elasticsearch as a Service offering. It can be used by the Docker community for search and discovery.	» Management/Orchestration
<b>Enterprise DC/OS (Mesosphere)</b> Mesosphere's DC/OS is a commercial version of the Mesos OS for managing data centers. It supports both Kubernetes and Docker.	» Management/Orchestration

Product/Project (Company or Supporting Org.)		Functionality Includes:
	<b><a href="#">Joyent Triton (Joyent)</a></b>	» Management/Orchestration A Containers as a Service offering, Triton comes with ContainerPilot, DataCenter and SmartOS functionality built-in.
Open Source	<b><a href="#">Kubernetes (Cloud Native Computing Foundation)</a></b>	» Management/Orchestration Kubernetes is an open source Docker orchestration tool. Google initially developed Kubernetes to help manage its own LXC containers. Stateful support is done through a new object called Pet Set. In addition, there are many networking and data volume plugins available.
	<b><a href="#">Nirmata (Nirmata)</a></b>	» Management/Orchestration A cloud-based platform for managing microservices.
	<b><a href="#">OpenShift Online (Red Hat)</a></b>	» Management/Orchestration Red Hat's hosted version of OpenShift, a container application platform that can span across multiple infrastructure footprints. It is built using Docker and Kubernetes technology.
	<b><a href="#">Platform9 Managed Kubernetes for Docker (Platform9)</a></b>	» Management/Orchestration Kubernetes offered as a managed service. Customers can utilize Platform9's single pane of glass, allowing users to orchestrate and manage containers alongside virtual machines. In other words, you can orchestrate VMs using OpenStack and/or Kubernetes.
	<b><a href="#">StackPointCloud (StackPointCloud)</a></b>	» Management/Orchestration Allows users to easily create, scale and manage Kubernetes clusters of any size with the cloud provider of their choice. Its goal is to be a universal control plane for Kubernetes clouds.
Open Source	<b><a href="#">SwarmKit (Docker)</a></b>	» Management/Orchestration <i>A toolkit for orchestrating distributed systems at any scale. It includes primitives for node discovery, raft-based consensus, task scheduling and more.</i>
	<b><a href="#">Tectonic (CoreOS)</a></b>	» Management/Orchestration A commercial distribution of combined Kubernetes and CoreOS stacks. Tectonic is a Kubernetes-based, end-to-end commercial container orchestration engine for the enterprise, with an added focus on security.

# MISCELLANEOUS

*Many tools are harder to group into a specific category. The majority of the remaining listings are for performance monitoring of applications and infrastructure.*

Product/Project (Company or Supporting Org.)	Functionality Includes:
<a href="#">Amazon CloudWatch (Amazon Web Services)</a> A monitoring service for AWS cloud resources and the applications that are run on AWS.	» Performance Monitoring
<a href="#">Bluemix App Monitoring and Analytics (IBM)</a> This component provides health and availability information about applications running in Bluemix. Features include log analytics, alerting and diagnostics for root cause analysis.	» Analytics
<a href="#">Bluemix Availability Monitoring (IBM)</a> Bluemix Availability Monitoring is integrated into the DevOps toolchain. It runs synthetic tests from locations around the world, around the clock, to proactively detect and fix performance issues before they impact users.	» Performance Monitoring
<a href="#">Datadog (Datadog)</a> Datadog is a monitoring and analytics service for IT operations and development teams. It has containerized agents that can monitor container environments.	» Performance Monitoring
<a href="#">Dataloop.IO (Dataloop Software)</a> A systems metrics monitoring service that is meant to be used by both DevOps and operations teams.	» Performance Monitoring
<a href="#">Elasticsearch (Elastic)</a> Elasticsearch is a search and analytics engine based on Lucene.	» Analytics
<a href="#">Infrastructure Monitoring and Reporting for Bluemix Infrastructure (IBM)</a> This is the cloud monitoring solution that previously came with the Softlayer offering. It is based on CA Nimsoft, which is now called CA Unified Infrastructure Management for Cloud Monitoring.	
<a href="#">Instana (Instana)</a> Application performance management to automatically monitor dynamic modern applications.	» Performance Monitoring
Open Source <a href="#">Nagios Core (Nagios Enterprises)</a> Nagios Core serves as the basic event scheduler, event processor, and alert manager for elements that are monitored. It features several APIs that are used to extend its capabilities to perform additional tasks, implemented as a daemon written in C (for performance), and designed to run natively on Linux/*nix systems.	» Performance Monitoring

Product/Project (Company or Supporting Org.)	Functionality Includes:
<b>Netsil (Netsil)</b>	» Performance Monitoring Netsil's Application Operations Center (AOC) helps site reliability engineers (SREs) and DevOps overcome complexity and collaboration issues in API and microservices-driven production applications. At the heart of the AOC is an auto-discovered service topology rendered from live interaction analysis. It allows practitioners to visualize their service dependencies and operational metrics, and work collaboratively across teams. This enables effective application monitoring, capacity planning, incident response, and deploy management.
<b>New Relic APM (New Relic)</b>	» Performance Monitoring Application performance monitoring is at the heart of New Relic's suite of products, which it starting to call Digital Intelligence Platform. Its agent-based approach is particularly good for troubleshooting code-related application performance issues.
<b>Opsclarity (Opsclarity)</b>	» Performance Monitoring Monitors the health of data frameworks and associated streaming applications.
<b>OverOps (OverOps)</b>	» Performance Monitoring Formerly known at Takipi, OverOps provides monitoring for Java-based applications.
<b>Project Shipped (Cisco)</b>	Cisco's Project Shipped is a model for deploying microservices to the cloud. To reduce environmental inconsistencies, Project Shipped also emulates the cloud environment on your development workstation. Vagrant, Consul and Cisco Microservices Infrastructure are components. It displays metrics such as total containers versus running containers, CPU and memory utilization, and network activity of Docker containers and the applications running in them.
Open Source <b>Sensu (Sensu)</b>	» Performance Monitoring A monitoring framework for self-hosted, centralized metrics services.
<b>Sensu Enterprise (Sensu)</b>	» Performance Monitoring Sensu Enterprise is the commercial version of Sensu, a monitoring framework for self-hosted, centralized metrics services.
Open Source <b>VAMP (Magnetic.io)</b>	Vamp stands for Very Awesome Microservices Platform. It helps developers build, deploy and manage microservices. Vamp's core features include a platform-agnostic microservices DSL, powerful A/B testing, canary releasing, autoscaling and an integrated metrics/event engine.
<b>X-Pack (Elastic)</b>	A full stack monitoring solution that adds additional functionality to the Elastic Stack (formerly ELK).
Open Source <b>Zabbix (Zabbix)</b>	» Performance Monitoring An open source, enterprise-class monitoring solution.

# DISCLOSURES

The following companies mentioned in this ebook are sponsors of The New Stack: Apcera, Aporeto, Arcadia Data, Bitnami, Cloud Foundry, CloudMunch, Cloudsoft, CNCF, Codenvy, CoreOS, DigitalOcean, GoDaddy, HPE, Intel, Iron.io, Macaw, Mesosphere, Nuage Networks, Red Hat OpenShift, Twistlock and Wercker.

